

Data Engineering and Machine Learning Pipeline

Phase II Project Report

ENGE707 Data Engineering and AI

Leo Cao (23211788), Fateh Bhular (23217987), Shawn Lee (23204035)

Auckland University of Technology

Table of Content

1. Introduction – Project Phase-I Summary.....	3
1.1 Problem Restatement.....	3
1.2 Dataset summary.....	3
1.3 Data-quality issues identified in Phase I.....	4
1.4 Justification for an LCS-based solution.....	5
2. Original LCS System on Raw Dataset.....	6
2.1 LCS variant selection (preferably eLCS) – justification.....	6
2.2 Confirmation that original LCS code is used without algorithmic modification.....	6
2.3 Execution on raw / minimally processed dataset.....	7
2.4 LCS parameter settings used.....	9
2.5 Baseline performance results.....	10
2.6 Minimal processing required for LCS compatibility.....	11
3. Data Preprocessing and Feature Engineering.....	13
3.1 Data leakage prevention measures.....	13
3.2 Data cleaning.....	14
3.3 Outlier investigation and handling.....	16
3.4 Feature engineering.....	17
3.5 Encoding categorical variables.....	18
3.6 Discretising & Encoding numerical variables.....	19
3.7 Class-imbalance handling.....	20
3.8 Expected impact of preprocessing on LCS performance.....	20
4. Improved LCS-Based System.....	21
4.1 Summary of changes vs. the original LCS system.....	21
4.2 Justification.....	23
4.3 Pseudocode of the improved system.....	24
4.4 Comparison against the original LCS baseline.....	26
5. Experiments.....	28
5.1 Validation strategy, Justification of using 80/20 split over k-fold CV.....	28
5.2 Evaluation of metrics.....	29
5.3 Statistical significance testing.....	31
6. Model Comparison.....	33
6.1 Choosing the three conventional models.....	33
6.2 Model Setup.....	34
6.3 Model Comparison.....	35
7. Interpretation of Results and Explainable AI.....	41
7.1 Inspection of the learned rule population.....	41
7.2 Interpretation of learned heart disease rules.....	42
7.3 Interpretability and trustworthiness.....	43
8. Discussion, Limitations, and Future Work.....	44
8.1 Reasons for performance improvement (or lack thereof).....	44
8.2 Practical significance of results beyond raw numbers.....	46
8.3 Limitations.....	47
8.4 Future work.....	49
Section 9: Contribution and Reflection.....	50
9.1 Contribution Statement.....	50
9.2 Individual Reflections.....	51
References.....	52

1. Introduction – Project Phase-I Summary

1.1 Problem Restatement

Heart disease is a major concern that currently affects 195,000 people in New Zealand alone. Approximately every 90 minutes, a New Zealander dies from heart disease, and it is the number one killer for women (*Latest Heart Disease Statistics*, n.d.). In this project, the dataset “281K Patient Dataset for Multi-Disease Prediction” from Kaggle is used. The goal is to find whether the lifestyle, demographics, and clinical measurements in the dataset have any relevance to predicting whether an individual has heart disease. A model with this prediction ability isn’t a full replacement for having an in-person cardiovascular risk assessment. Still, it could be a cost-effective way to check for an




Image: © bluringmedia/Getty Images

individual’s chance of having heart disease. The target column “heart_disease_target” is a binary classification created after cleaning and is based on the “sublabel” column, which contains the combination of diseases an instance has. If the instance from the sublabel contains “HT”, the value of the new column would be True; otherwise, False. This new column is created to be easily used to compare and train the machine learning model’s prediction capabilities.

1.2 Dataset summary

Unified Chronic Disease Dataset (281K Patients, 38 Features)

Overview

The **Unified Chronic Disease Dataset** is an integrated healthcare dataset developed by combining three publicly available datasets covering **diabetes, heart disease, and hypertension**. The integration was performed using a **Medallion Architecture (Bronze → Silver → Gold)** in Microsoft SQL Server to create a clean, standardized, and analytics-ready dataset for machine learning and healthcare research.

The final dataset contains **280,985 anonymized patient records** with **38 clinical and lifestyle features** and **no missing values**. Each record represents an individual patient’s demographic profile, health measurements, lifestyle factors, and disease indicators.

To support a wide range of predictive modeling tasks, the dataset provides both a **binary target (Normal or Abnormal)** and an **8-class disease-combination label**, enabling researchers to study individual diseases as well as co-morbidity patterns.

Figure 1.2.1: Dataset overview from Kaggle

The dataset is a combination of three datasets (diabetes, heart disease, and hypertension). There are 280,985 raw records, and 278,701 remain after removing duplicates; the 278,701 will be used instead of the raw amount. The original dataset has 39 columns, but the introduction of “heart_disease_target” gives 40 columns. The columns utilised during the analysis and processing stage are lifestyle, demographics, clinical, and disease indicators. No missing values were found during the dataset inspection.

The overall heart disease rate is 2.1% across all these sources, but individually the dataset shows varying heart disease rates: 20.06% in the heart disease dataset source, compared to 4.17% for the diabetes source, and 0% for the hypertension source.

1.3 Data-quality issues identified in Phase I

From the data inspection and data cleaning stage, multiple issues in the dataset were identified:

- 2,284 duplicate instances were removed.
- heart_disease has fractional values (from 0 to 1), which doesn't clearly explain whether an individual has heart disease, so a new column (heart_disease_target) was created to serve as the target (binary) instead of the heart_disease column.
- The columns sublabel, disease_flags, label, and heart_disease all contain data leakage information. This provides clues and partial answers to whether the individual has heart disease; thus, they are excluded from the prediction process.
- Some columns such as age/age_level/age_normalized and bmi/bmi_level carry the same information; this is to be reviewed before modelling.
- Imbalance in the distribution of instances with and without heart disease, since 97.9% of records do not have heart disease and 2.1% do.
- As described in the dataset summary, individual sources have different heart disease rates.

1.4 Justification for an LCS-based solution

From the exploratory data analysis in phase 1, the heat map demonstrated that there was a weak linear correlation with the heart_disease_target; e.g age is 0.08, BMI is 0.04, and triglycerides are 0.09. This suggests that a single feature alone does not have a strong correlation with the target. When just analysing BMI level and age group, a clear difference in heart disease rate was observed between different age groups and BMI levels; seniors who are also overweight displayed the highest heart disease rate of 6.53% compared to 2.1% of the overall heart disease rate. Adding gender results in a difference in heart disease rates of 8.18% for males and 4.99% for females who are seniors and overweight. The overall takeaway is that single features are weak as linear predictors, and that combining specific features together showed a greater heart disease rate.




Figure 1.4 Heart disease rate across combinations of age, BMI and gender from phase 1

The main reason why LCS is suitable for this problem is that it allows us to find a specific set of conditions that predicts the heart_disease_target well. In the exploratory data analysis, all possible combinations of features with the target could not be examined due to time constraints as well as the number of possible combinations of features in the dataset; therefore, LCS would help with automation. Another reason is that LCS involves a population to create if-then rules (classifiers). This could allow interpretability for each classifier, as the produced model aims to help predict heart disease. Experts in the field could analyse these attributes to see whether each classifier makes medical sense. However, there is a challenge: the dataset has a clear imbalance of instances with heart disease compared to those without; thus, the LCS machine learning model may favour the majority class. Further evaluation would be needed beyond just the accuracy.

2. Original LCS System on Raw Dataset

2.1 LCS variant selection (preferably eLCS) – justification

We used eLCS, also known as Education Learning Classifier System, as our LCS system for this project. eLCS is a variant of Learning Classifier systems that uses supervised learning and is used for classification problems. This project's goal is to successfully predict whether a patient has heart disease, which is represented by "heart_disease_target" with a value of True or False. A True/False value can be represented by binary 1 for True and 0 for False, making it a binary classification problem.

In our dataset, "281K patient dataset", there are both continuous and categorical variables. eLCS can work with both of these types of variables, which made it a suitable LCS for this project.

We learned in phase I that there were no relationships between individual features and the target attribute. However, after combining features together and then comparing them, the relationships were clearer. eLCS is capable of learning rules that include a combination of attributes, which is exactly what this project requires for a machine learning model.

2.2 Confirmation that original LCS code is used without algorithmic modification

```
setup(  
    name = 'scikit-eLCS',  
    packages = ['skeLCS'],  
    version = '1.2.4',  
    license='License :: OSI Approved :: GNU General Public License v3 (GPLv3)',  
    description = 'Educational Learning Classifier System',
```

Figure 2.2.1: eLCS version confirmation

Our original dataset was given to the original eLCS implementation. The version used was 1.2.4, and there were no modifications made to the source code or behaviour of the system. The setup remained untouched and was found in the “setup.py” file in the system files. External code was added outside the eLCS implementation package, but this was just to make the dataset ready for model training. The following code was added:

1. Loading the dataset
2. Creating the target attribute
3. Splitting the dataset into training and testing sets (80/20 split)
4. Predicting the results
5. Calculating all of the metrics from model training

2.3 Execution on raw / minimally processed dataset

```
Raw shape: (280985, 39)
Raw records: 280985
Raw columns: 39
Missing values: 0
Duplicate records: 2284
```

Figure 2.3.1: Raw Dataset

The raw dataset from “patient_data.csv” was used for the baseline eLCS model. As shown in Figure 2.3.1, there were 280,985 records and 39 columns. The raw dataset contained 2284 duplicate records, which were intentionally kept. The duplicate rows will be handled during Task 3 because the original eLCS model was meant to represent the raw dataset with none to very minimal preprocessing.

```
heart_disease_target
0    275081
1     5904
Name: count, dtype: int64
```

heart_disease_target	0	1
sublabel		
DI	31918	0
DI_HT	0	1398
DI_HT_HY	0	845
DI_HY	66473	0
HT	0	2595
HT_HY	0	1066
HY	69806	0
N	106884	0

Figure 2.3.2: Target Construction

The target attribute was not in the original dataset, so it was added after loading. The “sublabel” attribute was used to create it by checking if the value contained “HT”, which represented heart disease. The values including “HT” were mapped to class 1, while the values not including it were mapped to class 0.

The result of the target creation was that 275,081 records were mapped to class 0 (no heart disease), and 5,904 records were mapped to class 1 (heart disease). With these results, the positive class rate was approximately 2.10%.

```
Predictor shape: (280985, 36)
Target shape: (280985,)
Positive-class rate: 0.021011797782799792
```

Figure 2.3.3: Predictor Set

After the target attribute was added, there were a total of 40 columns. However, 4 of these were excluded from the predictor matrix due to having a one-to-one correlation with the target which would result in data leakage.

Therefore, the baseline experiment used 36 columns. We performed validation checks which confirmed the target attribute was binary and the dataset contained no missing values.

```
Training shape: (224788, 36)
Testing shape: (56197, 36)
Training positive rate: 0.021010908055590156
Testing positive rate: 0.021015356691638343
```

Figure 2.3.4: Train/Test sets

The dataset was split into training and testing sets with an 80/20 split. Stratification was used to preserve the class imbalance in both sets so that they could both be as close to the original dataset split as possible. The results shown in Figure 2.3.4 tell us that the training

set contained 224,788 rows (80% of rows from the dataset) and the testing set contained 56,197 rows (20% of the dataset). Both sets used 36 attributes because 4 columns from the original dataset had a one-to-one correlation with the target attribute. Both sets maintained the same positive rate of approximately 2.10%, which was the same rate in the raw dataset.

```
Starting seed 3..
Completed seed 3: balanced accuracy=0.5000, ROC-AUC=0.6925, PR-AUC=0.2241, positive predictions=0
Starting seed 6..
Completed seed 6: balanced accuracy=0.5000, ROC-AUC=0.7354, PR-AUC=0.1711, positive predictions=0
Starting seed 9..
Completed seed 9: balanced accuracy=0.5000, ROC-AUC=0.6086, PR-AUC=0.0782, positive predictions=0
Starting seed 12..
Completed seed 12: balanced accuracy=0.5000, ROC-AUC=0.6885, PR-AUC=0.1787, positive predictions=0
Starting seed 15..
Completed seed 15: balanced accuracy=0.5000, ROC-AUC=0.7252, PR-AUC=0.2188, positive predictions=0
Starting seed 18..
Completed seed 18: balanced accuracy=0.5000, ROC-AUC=0.6970, PR-AUC=0.2439, positive predictions=0
Starting seed 21..
Completed seed 21: balanced accuracy=0.5000, ROC-AUC=0.6830, PR-AUC=0.1565, positive predictions=0
Starting seed 24..
Completed seed 24: balanced accuracy=0.5000, ROC-AUC=0.7186, PR-AUC=0.1668, positive predictions=0
Starting seed 27..
Completed seed 27: balanced accuracy=0.5000, ROC-AUC=0.6563, PR-AUC=0.2024, positive predictions=0
Starting seed 30..
Completed seed 30: balanced accuracy=0.5000, ROC-AUC=0.6630, PR-AUC=0.1845, positive predictions=0
```

Figure 2.3.5: Results of the 10 baseline experiments

Figure 2.3.5 shows the results of the 10 repeated experiments using the original eLCS model. The seeds used were [3, 6, 9, 12, 15, 18, 21, 24, 27, 30], verifying that each experiment was unique. In every run, the seed affected both the shuffled split and the random state of the eLCS. Besides the seed, all of the hyperparameters were kept the same throughout all 10 experiments. We will use the same seeds in Part 4 for the improved eLCS model to make the comparison between data fair.

2.4 LCS parameter settings used

	Parameter	Value
0	N	1000
1	acc_sub	0.99
2	beta	0.2
3	chi	0.8
4	delta	0.1
5	discrete_attribute_limit	10
6	do_GA_subsumption	True
7	do_correct_set_subsumption	False
8	fitness_reduction	0.1
9	init_fit	0.01
10	learning_iterations	10000
11	match_for_missingness	False
12	mu	0.04
13	nu	5
14	p_spec	0.5
15	random_state	3
16	reboot_filename	None
17	selection_method	tournament
18	specified_attributes	[]
19	theta_GA	25
20	theta_del	20
21	theta_sel	0.5
22	theta_sub	20
23	track_accuracy_while_fit	False

Figure 2.4.1: Parameters List

When creating the baseline model, we used the seed as the random_state. So all **parameters were kept the same except for the seed** for each baseline experiment. The parameters in Figure 2.4.1 were retrieved using “model.get_params()”, which listed all of the parameters from the newly created eLCS object. Each of these affects how the experiment runs; for example, “learning_iterations” represents how many times eLCS performs learning iterations, “N” represents the maximum rule-population size, and “mu” represents the mutation probability. The rest of the explanations are in “Original_eLCS_Baseline.ipynb” in the project files.

```
# 10 reproducible seeds for the repeated experiments.
SEEDS = [3, 6, 9, 12, 15, 18, 21, 24, 27, 30]
✓ 0.0s
```

Figure 2.4.2: Seeds list

We used 10 different seeds (Figure 2.4.2) in **10 different and unrelated baseline experiments**. For each seed, a new eLCS model object was created, where the seed affected the train-test set split and the randomness of the experiment. The seed used for each experiment was recorded in order to make each run reproducible. As shown in Figure 2.3.5, the seed was recorded along with the experiment results for each run.

2.5 Baseline performance results

	Metric	mean	std	min	max
0	Accuracy	0.978985	1.170278e-16	0.978985	0.978985
1	Precision	0.000000	0.000000e+00	0.000000	0.000000
2	Recall	0.000000	0.000000e+00	0.000000	0.000000
3	F1	0.000000	0.000000e+00	0.000000	0.000000
4	Balanced accuracy	0.500000	0.000000e+00	0.500000	0.500000
5	ROC-AUC	0.686804	3.743284e-02	0.608623	0.735379
6	PR-AUC	0.182506	4.613999e-02	0.078235	0.243927
7	Positive predictions	0.000000	0.000000e+00	0.000000	0.000000
8	Training seconds	13.091345	1.869238e-01	12.765835	13.357953

Figure 2.5.1: Baseline eLCS results after runs

who actually did have a heart disease, the model was unable to identify any of them. This is extremely unsuccessful because the entire purpose for the model is to detect this target. Additionally, recall, and F1 score were all zero. Since there were no positive heart disease predictions, these scores cannot be calculated and are therefore zero. The balanced accuracy was 0.50 according to Figure 2.5.1. This means the model has chance-level performance, telling us that our model is the same as guessing. This is caused when the model only predicts one class and never the other.

The ROC-AUC and PR-AUC look at the probability scores. Some patients received higher scores, but because they weren't high enough, they weren't predicted as having heart disease. The 0.687 value of ROC-AUC means that the model can separate the two classes, but it is not strong enough to cause positive predictions. PR_AUC is more useful for this experiment because of the imbalanced dataset. The value is 0.183 for PR_AUC, which means that the model's probability scores ranked heart-disease cases better than random. Even though this checks the relationship between precision and recall, which were both zero, it only needs the scores, which is why PR_AUC isn't also zero.


	Confusion category	mean	std	min	max
0	True negatives	55016.0	0.0	55016.0	55016.0
1	False positives	0.0	0.0	0.0	0.0
2	False negatives	1181.0	0.0	1181.0	1181.0
3	True positives	0.0	0.0	0.0	0.0

Figure 2.5.2: Confusion Category Summary

The accuracy from the baseline experiment was approximately 97.9%. However, this is only because many records belong to class 0. If the model predicts every record as class 0, then the final accuracy will be 97.9%. Due to this, the value from the result does not give us the true model accuracy we are looking for. The precision of the baseline results is 0, which means for everyone in the data set

According to Figure 2.5.2, all experiments correctly predicted 55,016 class 0 cases, and incorrectly predicted 0 class 1 cases. There were no false positives or true positives, meaning the model never

predicted class one, and that no heart disease cases were detected. Zero false positives is usually a good thing, but in this case, it didn't predict heart disease at all, meaning that it missed all the positive cases, as shown by the false negatives.




```
Final training balanced accuracy: 0.5
Final training coverage: 1.0
```

Figure 2.5.3: Final accuracy and coverage

From the experiment results shown in Figure 2.5.3, the final balanced accuracy was 0.5, and the final coverage was 1.0. A training balanced accuracy of 0.5 tells us that the model could not distinguish between the two classes in the

training data. Because only class 0 was predicted and no class 1 predictions occurred, the equation for balanced accuracy returns 0.5. This means that the model only learned class 0, which was “no heart disease”. The training coverage of 1.0 tells us that all of the training samples matched at least one baseline eLCS rule. However, most of the rules predicted class 0, meaning that the rules were not effective enough to predict the minority class being class 1.

2.6 Minimal processing required for LCS compatibility



```
# all labels that include heart disease
HEART_DISEASE_LABELS = [
    "HT",
    "DI_HT",
    "HT_HY",
    "DI_HT_HY",
]

# copies the raw data and creates the binary target
df_baseline = df_raw.copy()

df_baseline["heart_disease_target"] = (
    df_baseline["sublabel"]
    .isin(HEART_DISEASE_LABELS)
    .astype(int)
)
```

Figure 2.6.1: Target construction

Before running the dataset through the baseline eLCS model, we had to make it compatible. Before, there was no target column for the model to make predictions into, so we created the heart_disease_target column as the target. The sublabel attribute was used to create the target value, which contained combined keys such as HT, DI, HY, and N. HT represented heart disease, so any records with HT in their sublabel were mapped to class 1 (heart disease), and the other records that

didn't contain HT were mapped to class 0 (no heart disease). The true/false values were converted to numeric 1 and 0 because the eLCS model requires numeric values. This was necessary in order for the model to use this dataset, and the original dataset was left untouched for the sake of backup.

```
# the target-related columns that cannot be predictors
EXCLUDED_COLUMNS = [
    "heart_disease_target",
    "sublabel",
    "disease_flags",
    "heart_disease",
]

# separates the predictor variables from the target
X = df_baseline.drop(columns=EXCLUDED_COLUMNS)
y = df_baseline["heart_disease_target"].astype(int)
```

Figure 2.6.2: Excluded Columns

```
# create a shuffled and stratified 80/20 train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.20,
    random_state=SEED,
    stratify=y,
    shuffle=True,
)
```

Figure 2.6.3: 80/20 split

Figure 2.6.2 shows how `heart_disease_target`, `sublabel`, `disease_flags`, and `heart_disease` were excluded from being used as predictors. `Heart_disease_target` was separated as the target, so it can't be the predictor. `Sublabel` was used to construct the target, so this also could not be used, as it would leak the answer. `Disease_flags` held information about disease, which could also leak the answer. Finally, `heart_disease` held information about the patient that the eLCS model could use to ignore every other value and guess whether the patient had heart disease or not. The 'label' column indicates if a patient is 'Normal' or 'Abnormal'. Although it does have a correlation with the target, it is very broad, and doesn't directly have a one to one correlation with heart disease as it could indicate different types of disease. Because of this, we decided to not remove this column from the baseline to keep the processing as minimal as possible.

Figure 2.6.3 shows how the 80/20 split was used to create the training and testing sets. Shuffling was used in order to reduce the influence of the original dataset's row order. The `test_size=20` makes sure that the test set is 20% of the records and the train set will be 80% of the records. Stratification was also used so that the 2.1% class 1 rate would remain in both sets; this is to keep both sets as similar to the original dataset as possible. Finally, a random seed was used so that this split could be reproducible.

```
# creates an encoder that can handle unseen test-set categories
encoder = OrdinalEncoder(
    handle_unknown="use_encoded_value",
    unknown_value=-1,
)

# copies the test/train sets before replacing categorical columns
X_train_encoded = X_train.copy()
X_test_encoded = X_test.copy()

# learns category mappings from the training data and applies them to both partitions.
X_train_encoded[categorical_columns] = encoder.fit_transform(
    X_train[categorical_columns]
)
X_test_encoded[categorical_columns] = encoder.transform(
    X_test[categorical_columns]
)
```

Figure 2.6.4: Category encoding

The categorical attributes in the dataset were converted into numerical values so that they could be processed through the eLCS model. Each categorical value was given a numerical code, and the encoder learned these mappings and applied it to the testing set. Using `fit_transform()` and `transform()` on the training and testing sets, respectively, prevented data leakage.

3. Data Preprocessing and Feature Engineering

3.1 Data leakage prevention measures

To prevent target leakage of `heart_disease_target`, there are multiple decisions made during the processing phase to address this issue. Two measures have been used to reduce the data leakage, Train/test leakage and target leakage.

```
Training shape: (222960, 39)
Testing shape: (55741, 39)

Training target distribution:
heart_disease_target
0    218247
1     4713
Name: count, dtype: int64

Testing target distribution:
heart_disease_target
0    54563
1     1178
Name: count, dtype: int64

Training positive rate:
2.1138%

Testing positive rate:
2.1133%
```

Figure 3.1.1: Train/Test leakage

The train/test split has been used before any other preprocessing measures. This is because splitting the data used for training and model testing should be divided. Models that are based on the same instances of data would cause the model to practically memorise the pattern instead of learning new patterns. A 80/20% split is used when 80% of the data is used for training while 20% is used for testing (The 80/20 split would be further explained in **section 5**). Additionally, splitting the testing and training data stops the testing data from influencing the preprocessing phase, e.g. calculating the median of glucose, as the testing data would change the median compared to just using the training data.

```
Removed target-leakage columns: ['sublabel', 'disease_flags', 'heart_disease', 'label']
Training shape: (222960, 35)
Testing shape: (55741, 35)
```

Figure 3.1.2: Data leakage

These columns label, sublabel, disease_flags, and heart_disease are removed in the processing stage because they have partial or full information about the target. In task 2, we did not remove the 'label' column because it was not a one to one correlation with the target, however it still does classify an individual as abnormal or normal with any type of disease, including heart disease so it will need to be addressed properly during preprocessing and removed. For the sublabel column, this was used to directly create the target because it contains the classification of diseases. The Disease_flags column contains the binary indicator of which disease the individual has. Lastly, the heart_disease column shows the probability of heart disease for a patient with similar traits and showed a strong correlation of 0.65 with the target; therefore, it cannot be used.

3.2 Data cleaning

In phase 1, the raw dataset has been cleaned for exploratory analysis (including missing values, duplicate records, inconsistent categories, and data types). Before going further, the dataset has been reviewed again to check if there are any issues after the splitting of the training and testing datasets. Invalid values were partially checked during phase 1 EDA; therefore, we will review them again.

We checked the missing values in the training and testing sets. No missing values were found. This is to ensure that the cleaning completed in phase 1 was inspected again.

```
Dataset shape: (278701, 40)
Missing values: 0
Exact duplicate records: 0

Target values:
[np.int8(0), np.int8(1)]

Target distribution:
heart_disease_target
0      272810
1       5891
Name: count, dtype: int64
```

3.2.1: Missing and duplicate records

When checking for the amount of duplication rows in the training and testing dataset after the removal of data leakage, showed that there are 25,812 instances in the training dataset and 2078 instances in the testing dataset that has at least 1 other instances with the same input feature values. The Figure 3.2.2 also shows the total number of repeated different combinations of input features that the training dataset has (11,285). Out of these 707 groups conflicting information about different instances with the same input features having different target values.

This is because matching predictor values do not inherently mean that the duplicate values are the same instance or is invalid, unlike during phase 1 where the identified duplicate rows were identical throughout the 39 columns which were removed in the cleaning process. The removal of all of these instances could result in discarding valid information as well as since our minority class (Heart disease) is already small, reducing these would worsen class imbalance.

Invalid values (Figure 3.2.3) were checked by making sure that the values are within valid boundaries, this include values that are negative and check that the values are within clinical measurements. Age, BMI, glucose, cholesterol, systolic_bp, diastolic_bp, and heart_rate all showed that the values are plausible.

```
composite_key:
  composite_key
Middle_Female_Overweight_Never_No 4441
Mature_Female_Overweight_Never_No 4302
Mature_Female_Obese_Never_No 4029
Middle_Female_Obese_Never_No 4011
Middle_Female_Normal_Never_No 3853
...
Young_Female_Underweight_Current_No 148
Young_Female_Underweight_Never_Yes 148
Young_Male_Underweight_Former_Yes 138
Young_Male_Underweight_Current_Yes 133
Young_Female_Underweight_Former_Yes 119
Name: count, Length: 192, dtype: int64

age_level:
  age_level
Mature 69654
Middle 67574
Senior 55874
Young 29858
Name: count, dtype: int64

gender:
...
hypertension 139929
diabetes 75162
heart 7869
Name: count, dtype: int64
Output is truncated. View as a scrollable element or open in a tool
```

3.2.4: Inconsistent categories

```
composite_key      str
age_level          str
gender            str
bmi_level          str
smoking           str
diabetes          str
age              int64
age_normalized     float64
bmi               float64
hypertension       float64
HbA1c_level       float64
glucose           float64
cholesterol        float64
sleep_hours       float64
triglycerides     float64
physical_activity  str
family_history     str
stress_level      str
low_hdl_cholesterol str
high_ldl_cholesterol str
blood_pressure    float64
high_blood_pressure str
sugar_consumption str
crp_level         float64
homocysteine_level float64
...
education_level   str
employment_status str
source_dataset    str
dtype: object
```

Figure 3.2.5: Data Types

All categorical columns were checked for inconsistent categories. Composite key, age_level, gender, bmi_level, smoking, diabetes, physical_activity, family_history, stress_level, low_hdl_cholesterol, high_ldl_cholesterol, high_blood_pressure, sugar_consumption, education_level, employment_status and source_dataset, all have consistent categorical values and formatting consistency.

The data types of each column are correct; therefore, no changes would be made (Figure 3.2.5).

3.3 Outlier investigation and handling

	feature	minimum	maximum	lqr_lower_bound	lqr_upper_bound	flagged_records
0	age	2.000000	89.000000	-22.000000	122.000000	0
1	hypertension	0.000000	1.000000	-0.325594	0.542657	6454
2	HbA1c_level	3.500000	9.000000	3.170101	9.083164	0
3	glucose	70.000000	300.000000	10.000000	250.000000	1624
4	cholesterol	150.000000	300.000000	138.000000	306.000000	0
5	sleep_hours	4.000000	10.000000	3.950000	9.950000	1216
6	triglycerides	50.000000	400.000000	-42.625000	412.375000	0
7	blood_pressure	120.000000	180.000000	146.284639	153.382711	24303
8	crp_level	0.003647	14.996625	6.599316	8.454110	24155
9	homocysteine_level	5.000236	19.999037	11.690145	13.283792	24774
10	systolic_bp	90.000000	179.000000	96.500000	172.500000	21854
11	diastolic_bp	60.000000	119.000000	63.500000	115.500000	18579
12	alcohol_intake	0.000000	30.000000	2.750000	27.150000	26112
13	salt_intake	2.000000	15.000000	3.300000	13.700000	26854

In phase 1, BMI was found to have a considerable number of outliers, but is proven that the outliers shouldn't be

transformed or removed. The other columns were discussed at a surface level, based on the boxplot, and showed no isolated instances. To further investigate this, the Interquartile range method was performed on the numerical columns. Triglyceride, cholesterol, and age showed no outliers, but the rest of the columns' results displayed incredibly high values, e.g 18579 outliers for diastolic_bp as well as other columns. A histogram is created to show the distribution of instances across each column. Most of the columns (systolic_bp, diastolic_bp, heart_rate, hdl, ldl, glucose) showed a high peak in the middle




Figure 3.3.2: Outlier investigation

and consistent values across the whole range. The IQR method identified many statistical outliers, however this does not automatically make them invalid. Inspecting the histograms and range values, they did not provide enough evidence to justify that these identified outliers would be treated as valid instances. Thus, no removal, transformation, or capping is performed.

3.4 Feature engineering

```
Removed redundant and provenance columns: ['composite_key', 'source_dataset', 'age_level', 'age_normalized', 'bmi']
Training shape: (222960, 30)
Testing shape: (55741, 30)
```

Figure 3.4.1: Removal of redundant and provenance variables

In the current dataset up to this point of processing, multiple columns are irrelevant or provide no value towards predicting the target, these include:

Composite key (removed): A composite key is just the combination of variables (age_level, gender, bmi_level, smoking, and diabetes); therefore, this does not provide any value, as it is simply grouping the variables for instance. So removing this would not cause any information to be lost and is just an unnecessary column to provide to the model.

Age_level and age_normalized (removed): Age_level and age_normalized provide similar information to age, but age_level is a categorical level and age_normalized is the linear scaling of age and showed a 1.00 correlation with age during EDA. These add no additional information about age, as age could be discretised independently for eLCS, thus being a more favourable column to remain in the training dataset.

bmi (removed): The bmi column is removed because the bmi_level column follows the WHO standard, which is the standard classifier for BMI. The bmi column does allow independent interpretation of discretisation, but since bmi_level is already in a suitable standard, there isn't any reason not to use bmi_level instead of bmi.

Source_dataset (removed): The source dataset has been removed because the dataset "281K Patient Dataset for Multi-Disease Prediction" is from the combination of three separate dataset. Each dataset varies in the proportions of instances with heart disease, therefore the model could use this to learn the pattern of which dataset is more likely to have heart disease resulting in its influence on its prediction.

3.5 Encoding categorical variables

	gender	diabetes	family_history	low_bld_cholesterol	high_bld_cholesterol	high_blood_pressure
226737	1	1	0	0	0	0
240229	0	0	1	0	1	0
263800	1	0	0	0	0	1
102871	1	1	1	0	1	0
209771	0	1	0	0	0	1

	bmi_level	physical_activity	stress_level	sugar_consumption	education_level
226737	3	1	0	1	1
240229	0	0	1	2	0
263800	2	0	1	2	0
102871	3	1	2	2	2
209771	0	1	1	1	0

	smoking_Current	smoking_Former	smoking_Never	employment_status_Employed	employment_status_Retired	employment_status_Unemployed
226737	0	0	1	0	0	1
240229	0	1	0	0	0	1
263800	1	0	0	0	1	0
102871	1	0	0	0	0	1
209771	0	0	1	0	0	1

Figure 3.5.1: Encoded columns

All the categorical variables have been encoded to numerical values because eLCS could not accept values in strings. For example, gender has been changed from "Female" and "Male" to 0 and 1.

Multiple types of encoding have been used for the categorical variables since some variables required different encoding methods.

Binary Encoded

For these columns (gender, diabetes, family_history, low_hdl_cholesterol, high_ldl_cholesterol, and high_blood_pressure), binary encoding is used because these columns only have 2 possible outputs.

Ordinal Encoded

For these columns (bmi_level, physical_activity, stress_level, sugar_consumption, "education_level), ordinal encoding is used because, for example, for stress_level, there is an order of the output values "low", "medium," and "high"; therefore, the encoding converts it into numerical values (1 to 3).

One-Hot Encoded

For the rest of the categorical columns ("smoking", "employment_status"), one-hot encoding is used since they provide no order and cannot be represented in binary; therefore, new columns would be created to represent each possible outcome, as shown in Figure 3.51 .

3.6 Discretising & Encoding numerical variables




Figure 3.6.1 Discretise and Encode the numerical attribute

Discretization has been used for all the attributes in Figure 3.6.1, changing the numerical dataset to categorical values because it simplifies feature representation as well as making it easier to interpret the rule. Even though eLCS is able to utilize the ranges on a classifier, removing the number of grouping would reduce the number of possible distinct values thus may make the rule learning much easier. The downside is the loss of detail in the data which would make the rule more generalized and could affect

eLCS predictive performance. To determine the categorical values, the columns that have a clinical standard include a fixed-boundary, the boundaries set are low, medium or high each with their own range for each categorical value. For example, for glucose, the normal range would be below or at 100mg/dL, the prediabetes level from 100 to 126mg/dL and diabetes at 126mg/dL and above.

The columns (hypertension, alcohol_intake and salt_intake) that don't have a set standard for measuring different levels are based on training derived quantile based boundaries. This is to roughly split training instances across 3 categorical levels (3 bins) based on the quantile percentage that each instance resides in.

These attributes are then encoded using ordinal encoding as the categorical values are arranged in order. The encoding is applied on both the training and testing dataset as well as the original attributes are removed and have been replaced by the created bins.

3.7 Class-imbalance handling

In the EDA, the target variable is highly imbalanced across the dataset, with 2.1% of instances with heart disease and 97.9% of instances without heart disease of the total 278,701 instances. The baseline eLCS prediction results showed that the final balanced accuracy test is 0.5. This balanced accuracy occurred because the baseline model practically predicted only the class which has no disease, and failed to identify any with a heart disease. This shows that the model does not provide any useful balanced predictions between the two classes.

The imbalance would be handled after receiving the result from the prediction preprocessed dataset (*Task 4*) in eLCS. This is because it would be valuable insight into how much class imbalance could affect the overall prediction of the model when the imbalance is handled.



Figure 4. Distribution of the heart disease target in the cleaned dataset from phase 1

3.8 Expected impact of preprocessing on LCS performance

The overall preprocessing phase of the dataset is expected to improve the eLCS ability to process and performance, these improvements include:

- **Improved Computational cost / quicker prediction:** In the feature selection process, multiple redundant columns were removed, therefore there are less combinations of if and then rules.
- **More compact rule set:** Applying the discretization lowers the amount of classifiers since there will be less classifiers that are tied to continuous values which also increase readability for validation and inspection of the rule.
- **Improved accuracy of rules:** Using the appropriate encoding types would reduce the number of misleading information when just applying a general encoding type to all columns which causes the eLCS to learn inaccurate patterns.
- **Increase trustworthiness of prediction:** This is because the columns that are removed that provide information about the target information would influence eLCS ability to learn and accurately predict the target without genuinely learning from the other features.

4. Improved LCS-Based System

4.1 Summary of changes vs. the original LCS system




Figure 4.1.1: Training distributions before and after undersampling

```
Full training shape: (222960, 34)
Testing shape: (55741, 34)

Full training target distribution:
heart_disease_target
0    218247
1     4713
Name: count, dtype: int64

Testing target distribution:
heart_disease_target
0     54563
1      1178
Name: count, dtype: int64

Undersampled training shape: (29713, 34)

Undersampled target distribution:
heart_disease_target
0     25000
1      4713
Name: count, dtype: int64

Negative-to-positive ratio: 5.30:1
```

Figure 4.1.2: Training statistics comparison

This section focused on improving the predictions of the eLCS model by using undersampling and tuning the hyperparameters. The full preprocessed data information is displayed in Figure 4.1.1. There were a total of 222,960 records in

the training set, where 218,247 of them were class 0, and 4,713 of them were class 1. Class 0 significantly outnumbers class 1 in the total number of records, which results in a class imbalance problem. This directly correlates with another issue, which was the model hyperparameters not fitting the dataset. The default population and learning_iteration values were not making use of all the data provided and didn't give the model enough time for training.

To address the class imbalance issue, sampling was used to randomly select 25,000 class 0 records to keep in the training set. The rest of the class 0 records were removed from the training set, resulting in a total training set size of 29,713 records. This resulted in a better ratio of 5.3 class 0 records for every

class 1 record. Before changing the hyperparameters, the new undersampled training set was used to train the model, just to see what different undersampling did by itself.

Seed	Accuracy	Precision	Recall	F1	Balanced accuracy	ROC-AUC	PR-AUC	True negatives	False positives	False negatives	True positives	Positive predictions	Training seconds
0	3	0.978723	0.100000	0.000849	0.001684	0.500342	0.601036	0.051937	54554	9	1177	1	11.185162
1	8	0.978867	0.000000	0.000000	0.000000	0.500000	0.579673	0.031664	54563	0	1178	0	11.300214
2	9	0.978867	0.000000	0.000000	0.000000	0.500000	0.489641	0.022186	54563	0	1178	0	11.221245
3	12	0.978867	0.000000	0.000000	0.000000	0.500000	0.551384	0.040587	54563	0	1178	0	12.379079
4	15	0.978867	0.000000	0.000000	0.000000	0.500000	0.599107	0.05318	54563	0	1178	0	11.700800
5	18	0.978867	0.000000	0.000000	0.000000	0.500000	0.437362	0.025779	54563	0	1178	0	11.522431
6	21	0.978867	0.000000	0.000000	0.000000	0.500000	0.583367	0.046269	54563	0	1178	0	11.865966
7	24	0.978597	0.222222	0.006093	0.009959	0.502354	0.571226	0.040622	54542	21	1172	8	11.656360
8	27	0.978867	0.000000	0.000000	0.000000	0.500000	0.570092	0.035618	54563	0	1178	0	11.256830
9	30	0.978867	0.000000	0.000000	0.000000	0.500000	0.531064	0.026285	54563	0	1178	0	11.527865

Seed	Accuracy	Precision	Recall	F1	Balanced accuracy	ROC-AUC	PR-AUC	True negatives	False positives	False negatives	True positives	Positive predictions	Training seconds
0	3	0.977593	0.211382	0.022071	0.039969	0.510147	0.795535	0.112511	54466	97	1152	26	5.231056
1	6	0.971888	0.194682	0.105263	0.136639	0.547931	0.809791	0.113939	54050	513	1054	124	5.396431
2	9	0.976373	0.216327	0.044992	0.074491	0.520736	0.754192	0.099405	54371	192	1125	53	6.183733
3	12	0.967241	0.183594	0.159593	0.170754	0.572135	0.734273	0.104268	53727	836	990	188	5.223253
4	15	0.977073	0.191358	0.026316	0.046269	0.511957	0.755534	0.084237	54432	131	1147	31	5.083849
5	18	0.974794	0.171014	0.050085	0.077479	0.522422	0.806915	0.112479	54277	286	1119	59	5.426928
6	21	0.978803	0.239819	0.044992	0.075768	0.520956	0.753362	0.089727	54395	168	1125	53	5.109017
7	24	0.978813	0.200000	0.000849	0.001691	0.500388	0.752372	0.103252	54559	4	1177	1	5.391051
8	27	0.978526	0.256410	0.008489	0.016434	0.503979	0.795370	0.113459	54634	29	1168	10	5.572949
9	30	0.978463	0.203540	0.039049	0.065527	0.517875	0.735410	0.082736	54383	180	1132	46	5.360582

Figure 4.1.2: Preprocessing vs Undersampling results

Figure 4.1.2 shows the difference between using the fully preprocessed and undersampled training set. The undersampled training experiment produced much better results, because the default eLCS model was able to predict many more true positives in all 10 experiments. False negatives are much more critical in healthcare, because it means deadly diseases go unnoticed until it's too late. Fewer false negatives and more true positives in the undersampled experiments were a sign that it produced better results than the preprocessed experiments.

Configuration	Training data	N	Learning iterations
0 Default eLCS	Full or undersampled	1000	10000
1 Improved eLCS	Undersampled	8000	50000

Figure 4.1.3: Improving the eLCS model

After checking that undersampling succeeded, we had to change the eLCS hyperparameters so that it fit the needs of the undersampled

training set. N represents the population of the rules, which by default is 1,000 in eLCS.


This is very little for our dataset and restricts the model's learning capacity. Learning.iterations is how many records the model will go through in an experiment, and 10,000 iterations means that it will go through 10,000 records in each experiment. This is a significant limitation for our training set, because it means it will only go through 1/3 of the training set in each experiment (skipping most of the training

set). To overcome this limitation, the learning iterations were increased from 10,000 to 50,000 so that it would go over every record at least once in each experiment. Increasing it could improve the experiment results, but would also increase computational costs. Everyday laptops aren't made for machine learning, so we decided to limit the learning iterations to 50,000. There was no modification of the eLCS source code, only the hyperparameters.

There were many variables that did not change throughout the different experiments. The first one was the test set, which was kept the same with 55,741 records (54,563 class 0 records and 1,178 class 1 records). The same seeds were used throughout every 10 repeated experiments in tasks 2 - 4 to keep the experiments reproducible and not introduce any biases. Separating these changes into 3 parts allowed us to understand how each change affected the final predictions made by the eLCS model.

4.2 Justification

Sampling was used to solve the class imbalance issues in our dataset. The full training dataset contained approximately 46 negative records for every single positive record. This led to the eLCS systems predominantly learning only the class 0 rules. At first, the high accuracy of 97.9% was interpreted as good; however, this is misleading because predicting class 0 for everything gives a 97.9% accuracy. Reducing the class 0 records gave class 1 records more influence during the rule discovery process. Class 1 records were all retained because we could not afford to lose any real heart-disease examples. The other solution was to generate new class 1 records, but not only would artificial records be bad for the model training, but it would also increase the computational cost.



Changing the N hyperparameter to 8,000 allowed the model's population to contain more candidate rules. The purpose of increasing the population was to give the model more capacity to be able to retain classifiers/patterns that would otherwise be competing with each other in a smaller

population. Changing learning iterations to 50,000 gives rules more opportunities during the experiments to be matched together, evaluated, reproduced, and refined. Increasing the iterations also means that the model is able to properly go through and learn from all the training instances at least once. While changing the N and learning iterations hyperparameters resulted in better results across ROC-AUC, PR-AUC and precision, it also had a higher computational cost.

4.3 Pseudocode of the improved system

```
Split the cleaned data into stratified training and testing sets

Fit preprocessing transformations on the training data
Apply the fitted transformations to the test data

Save the fully preprocessed training data

positive_records = all class 1 training records
negative_records = randomly selected sample of 25,000 class 0 records
undersampled_training = shuffle(positive_records + negative_records)

For each seed in [3, 6, 9, 12, 15, 18, 21, 24, 27, 30]
    Model A:
        Train default eLCS model on fully preprocessed training data
    Model B:
        Train default eLCS model on undersampled training data
    Model C:
        Train improved eLCS (N=8000, Iter=50000) on undersampled
        training data.

Evaluate all models on the same test set (remains unchanged)
Store the metrics, confusion matrix, and training time
```

Figure 4.3.1: Explanation of the improved system in Pseudocode

Figure 4.3.1 describes how the improved eLCS system works in a simplified way. It starts by using the phase 1 dataset and divides it into stratified training and testing sets with an 80/20 split. The preprocessing is learned from the training data and then applied to the testing set, while preventing test-data leakage. All class 1 records are kept, while 25,000 randomly selected class 0 records are retained. There are three different configurations of models that allow each change to be inspected

individually. All models use the same test set, which was kept near the beginning of the code, and are compared, with the results being stored in files.




Figure 4.3.2: Improved eLCS experimental pipeline

This flowchart in Figure 4.3.2 gives a visual overview of how the data moves through the improved eLCS pipeline. The split that happens after leakage-free preprocessing demonstrates how data is split into two groups, which are the “full-training” branch and the “undersampled” branch. The left branch works with the data that hasn’t been through imbalance treatment. The right branch works with data that goes through imbalance treatment and through both the baseline and improved eLCS models. The convergence of both branches visually demonstrates how they both use the same testing set so that no model has an advantage over the other.

This entire pipeline is reproducible due to all information being recorded and saved, like the 80/20 split information, the sampling process,

and the same 10 eLCS seeds that are used during the repeated tests. This flow chart is a visual version of the pseudocode from Figure 4.3.1, which explains how the system works in a simplified way.

4.4 Comparison against the original LCS baseline

This section discusses the comparison between the baseline eLCS model trained in Section 2 and the newly improved eLCS model. The **baseline eLCS model** used the minimally processed dataset, which included modifications just to be compatible with the model. This experiment did not address class imbalance, and used a population of 1,000 rules, 10,000 learning iterations, and seed-specific 80/20 train-test sets. The **improved eLCS model** used the cleaned and preprocessed dataset where sampling was applied to the training set. The sampling retained 4,713 class 1 records and 25,000 randomly selected class 0 records, with a total of 29,713 records in the training set. The population of rules was increased from 1,000 to 8,000, and the learning iterations from 10,000 to 50,000. The improved eLCS model also used a different testing set from the baseline model, which was the test set created in task 3. This is not a perfectly controlled experiment as several factors were changed, so the differences in results can be from preprocessing, undersampling, hyperparameter tuning, or all of these together.





Figure 4.4.1: Comparison of baseline vs improved

The values of both models in Figure 4.4.1 are the results of the model differences. 5 out of 6 metrics (*explained in section 5*) were higher in the improved eLCS, indicating that it performed better than the baseline model. **See section 6 for in depth evaluation.**



The baseline eLCS model had a mean training time of 38.8 seconds while the improved model had a mean training time of 224.4 seconds. This means the improved model was approximately 5.8 times slower. The main reason why the computational cost increased so much was because of the hyperparameter tuning. The rule population increased from 1,000 to 8,000 and the learning iterations from 10,000 to 50,000. Undersampling did reduce the number of records, but the larger population and more learning iterations had a bigger effect on the computational cost. The extra cost allowed the model to produce better ROC-AUC scores meaning the model could make correct positive predictions, but it only made minor improvements to recall, F1 and balanced accuracy. Without recall increasing, the computational cost can not be stated as being worth it.

5. Experiments

5.1 Validation strategy, Justification of using 80/20 split over k-fold CV

We went with an 80/20 split, which split the dataset into a training set with 80% of the data and a testing set with 20% of the data. In the original dataset, approximately 2.1% of patients had heart disease. This is a very low number and introduces class imbalance to the dataset. To maintain this environment, stratification was used to keep




Figure 5.1 Class Distribution After Stratified 80/20 Split

the same 2.1% proportion, making the training and testing sets similar to the raw dataset. A seed was recorded and used for the split so that the 80/20 split could be reproducible, just in case we wanted to re-run the split again with the same results. The exact same seed was used in the baseline model to make sure the comparison between the baseline eLCS model and the improved model was fair. The experiment was repeated using 10 recorded random seeds so that the results were not based on a single train-test split.

The 80/20 split was a better approach than k-fold cross-validation, mainly because of the size of the dataset. And the computational requirements. K-fold CV requires the model to be trained multiple times, with K representing the number of times the model will be trained. The eLCS model keeps on creating, evaluating, and evolving rules in each training, so this whole process happening repeatedly is very computationally costly. The dataset contains over 200k rows, which is more than enough for an 80/20 split to be effective for machine learning. If the dataset were smaller, then it would be less reliable because the training set would not contain enough data to train the model.

lower-upper bound, prediction array, crossover mutation, 7-8k population, 50,000 iterations, 25,000 majority class instances, match set, covering

5.2 Evaluation of metrics

Since our data set is heavily unbalanced, with only around 2.1% of the records being classified as having a heart disease. Using accuracy alone is not a good enough metric to evaluate the performance of our model. The main metrics that were used in this project are **recall, PR-AUC, and balanced accuracy**; the other metrics such as **precision, F1-score, ROC-AUC, accuracy, and the confusion matrix** are instead used as supporting measures alongside the main evaluation metrics.

Recall measures the percentage of actual positive cases that the model correctly identifies. In the context of this project, the recall tells us that out of all patients who have a heart disease, how many of them were correctly identified. Recall is especially important in our context because the main objective of our model is to identify true positive cases. Recall has been given more importance for this project because the main objective is to identify positive heart disease cases. A false positive is when a patient

is classified as having a heart disease even though they don't, however this scenario is more desirable than if the model were to miss an actual positive case (false negative)

Precision measures the percentage of all the positive cases identified by the model being actual positives. In the context of this project, the precision tells us how many of the patients that were identified with heart diseases genuinely do have a heart disease. Precision is important because a model that predicts too many false positives would be less practicable and lose credibility. However we are treating precision not as important as recall for this project because being able to identify more actual heart disease cases is more important, even when with the compromise of more false positives.

The **F1 score** combines the recall and precision into one measurement. This is useful for this project because it has the desired effect of evaluating if the model can correctly identify positive cases, however at the same time considering false positives. The F1 score is less affected by the large imbalance of negative cases which makes this a much more suitable metric for this highly imbalanced classification problem. However, recall is still placed higher for this project because F1 gives precision and recall equal importance, however this project aims to identify as many heart disease cases as possible.

The **balanced accuracy** shows the average performance/recall between the two classes. This is useful for our context because each class will contribute an equal amount to the metric which is good against highly imbalanced datasets. Comparing this with normal accuracy, it can still remain very skewed due to this imbalance while balanced accuracy skews less.

The **PR-AUC** measures how well the model is able to identify positive cases while still limiting the amount of false positives. This is especially important for this project because identifying the positive cases (heart disease) is the main objective, however they are rare, only contributing around 2.1% of the records.

The **ROC-AUC** measures how well the model is actually able to separate the positive cases from the negative cases. This is useful for evaluating the overall ability that model has of ranking classes.

However, this is not as important as PR-AUC because a high ROC-AUC does not necessarily mean more true positive cases are actually identified.

Accuracy measures the percentage that all the predictions made by the model are correct. The accuracy has been reported for each of the models we ran, however it is not used as the main evaluation metric because our dataset is heavily imbalanced. Since 97.9% of the records fall into class 0 (no heart disease), a model which just predicts class 0 for every single instance could achieve an accuracy of 97.9% but still fails to identify even a single heart disease case. This is why for our case, the accuracy is instead interpreted alongside other evaluation metrics that are better designed for class unbalances.

The **confusion matrix** shows how many predictions were correct and incorrect for each class. It shows the number of true positives, true negatives, false positives, and false negatives. This is useful for our context because it directly tells us how many heart disease cases were detected by the model, and how many it missed.

Improved eLCS Prediction Outcomes

Actual class	Predicted class	
	Class 0 No heart disease	Class 1 Heart disease
Class 0 No heart disease	True Negative 99.94% Mean: 54,529.0	False Positive 0.06% Mean: 34.0
Class 1 Heart disease	False Negative 99.07% Mean: 1,167.0	True Positive 0.93% Mean: 11.0

After evaluation of all the metrics, the **recall**, **PR-AUC**, and **balanced accuracy** are the main metrics used to determine how well our model performs in this project. The recall is able to show how much heart disease cases in the dataset have been recorded by our model, the PR-AUC takes this a step further and measures how well the model identifies positive cases while still limiting the amount of false positives, and lastly the balanced accuracy is a good measure to see how well the model performs overall on the two classes without too much bias from an unbalanced dataset. The remaining metrics are still useful, however they are more credible when interpreted alongside the main metrics.

5.3 Statistical significance testing

A paired t-test was used to see whether or not the improved changes from section 4 were statistically significant across the 10 repeated runs. The test results were paired together by random seed, so seed 3 was compared with seed 3, seed 6 with seed 6, and so on.

The null hypothesis is that there is no improvement between the two changes done to the model, and that any differences are due to variance between seeds. The significance level used was $p < 0.05$, and the main evaluation metrics chosen to be tested were recall, PR-AUC, and balanced accuracy as mentioned in 5.2.

	Change tested	Metric	Before	After	t-statistic	p-value	Result
0	Add undersampling	Recall	0.0006	0.0502	-3.223	0.01043	Significant increase
1	Add undersampling	PR-AUC	0.0376	0.1016	-11.196	1.387e-06	Significant increase
2	Add undersampling	Balanced accuracy	0.5003	0.5229	-3.248	0.01003	Significant increase
3	Increase population size + training iterations	Recall	0.0502	0.0093	2.542	0.0316	Significant decrease
4	Increase population size + training iterations	PR-AUC	0.1016	0.1783	-20.302	7.953e-09	Significant increase
5	Increase population size + training iterations	Balanced accuracy	0.5229	0.5044	2.542	0.0316	Significant decrease

Figure 5.3.1: Statistical significance results for the Task 4 eLCS model comparisons

Looking at the results for the first comparison, adding undersampling was able to increase the recall from 0.0006 to 0.0502, PR-AUC from 0.0376 to 0.1016, and balanced accuracy from 0.5003 to 0.5229. Furthermore, all three p values are under 0.05 which means that the increases were statically significant. These show that undersampling genuinely improved the eLCS across the three metrics, and they are not caused by random variance.

The addition of changing the eLCS hyperparameters gave mixed results. The population size and training iterations improved the PR-AUC, which went up from 0.1016 to 0.1783. But it had a large impact

on the recall of the model, which went down from 0.0502 to 0.0093 and the balanced accuracy down from 0.5229 to 0.5044, the decreases also statistically significant.

This testing shows that undersampling gave the most successful improvements to the model, and the increase in population size and number of training iterations resulted in an increase of PR-AUC, however it came with the cost in recall and balanced accuracy.

6. Model Comparison

6.1 Choosing the three conventional models

Logistic regression, decision tree, and random forest would be used for the comparison of the eLCS model prediction results. These three chosen models are all classification models, which allow for a valid comparison because the heart_disease_target is a classification problem. These represent different ways of predicting the target outcome, which would be helpful to compare their result metrics, computational cost, and interpretability.

Logistic Regression

For the logistic regression model, it utilises weights to form predictions. The weight change is determined by the loss function, which measures how similar the predicted probability is to the actual target value during training. Once the loss value is found using the loss function, the weight would need to be adjusted so that the loss value decreases. Logistic regression is used for comparison since this is a simple classification model which would provide a good benchmark to see if a simpler model is able to achieve a comparable performance to eLCS.

Decision Tree

A decision tree is similar to eLCS, where both utilize if-then rules, but how the rules are formed is different. Decision Trees generate rules by checking which split (condition) gets the purest group (The most dominant class). The group of instances would only stop once it reaches a leaf node (the last node of the tree branch) or a specified limit is introduced. The majority class of that node would be the prediction made. Decision trees have similar Interpretivity towards eLCS; therefore, it has been chosen for the comparison.

Random Forest

Random forest utilizes multiple decision tree predictions as votes, with the prediction with the most common vote being represented as the collective prediction. Each decision tree is trained on a different

subset of data; thus, their error could be different, so mistakes in prediction would likely be outvoted. This helps with lowering the instability and overfitting compared to using a single decision tree. Random forest is widely regarded as a strong benchmark model but is hard to interpret. If eLCS can match its predictive performance, it would provide comparable accuracy with greater interpretability.

6.2 Model Setup

The eLCS results of the baseline eLCS, the preprocessed eLCS and the improved eLCS (that contains changes in the hyperparameters and undersampling) are compared with the newly introduced conventional models (logistic regression, decision tree and random forest). These models were tested using the same 80/20 split performed in task 2. These models were also given the same 10 seeds to ensure that the comparison is as fair as possible as well as the same split dataset (the preprocessed dataset with undersampling). The 10 seeds would then be averaged out and find the mean and standard deviation. The parameters of the Logistic regression, decision tree and forest setting were mostly default values, but 2 hyperparameters were changed. The max_iter of the logistic regression were changed because the default was 100, setting it to 2000 ensures as a precaution that the convergence is reached. The other changed hyperparameter is the n_jobs in random forest, n_jobs were used to speed up the process of the training so it does not change the outcome of the model.

Even though the baseline and preprocessed eLCS weren't the fairest comparison as they used the dataset with class imbalance, they provide a good reference point of how much the improved eLCS made. The fairest test is against the improved eLCS and the three conventional models trained on the same preprocessed, undersampled data as they have the most similar conditions.

```
# Some seeds as Task 4.
SEEDS = [3, 6, 9, 12, 15, 18, 21, 24, 27, 30]

# Settings remain constant across runs.
LOGISTIC_SETTINGS = {
    "solver": "lbfgs",
    "max_iter": 2000,
    "C": 1.0,
    "class_weight": None,
}

TREE_SETTINGS = {
    "criterion": "gini",
    "max_depth": None,
    "min_samples_split": 2,
    "min_samples_leaf": 1,
    "class_weight": None,
}

FOREST_SETTINGS = {
    "n_estimators": 100,
    "criterion": "gini",
    "max_depth": None,
    "min_samples_split": 2,
    "min_samples_leaf": 1,
    "class_weight": None,
    "n_jobs": -1,
}
```

Figure 6.2.1: 3 conventional model specification

6.3 Model Comparison




Figure 6.3.1 : Bar Graph metric comparison

The baseline eLCS, preprocessed eLCS and the improved eLCS (undersampling & hyperparameter change) have very similar balanced accuracy, recall and F1. The other three models have a greater balanced accuracy, recall and F1 score. Random forest has the highest measurements across these three measurements of approx 0.788 balanced accuracy, approx 0.644 for recall and approx 0.267 for F1. The improved eLCS showed a very minor increase in F1 from the baseline and preprocessed configurations of only 0.017 and 0.016, respectively. Precision showed the improved eLCS has the highest value at 0.260 but also has the largest standard deviation of the different seed results.

The improved eLCS remained competitive with Random Forest on ROC-AUC and PR-AUC. Both models achieved ROC-AUC values of around 0.926 and 0.927, respectively, while improved eLCS achieved a PR-AUC of approximately 0.178 compared with approximately 0.168 for Random Forest. The baseline PR-AUC was also similar to the improved eLCS and Random Forest but had a wider standard deviation.

Overall the improved eLCS metrics are competitive with the other 3 conventional models such as ROC-AUC and PR-AUC but the balanced accuracy, recall and F1 are clearly overshadowed by the performance of the logistic regression, decision tree and especially the random forest. These values would be further evaluated in task 7.




Figure 6.3.2 : Confusion matrices of all models

In the confusion matrices, we can see that for all the eLCS configuration models (baseline, preprocessed and improved) almost all instances were predicted as negative. The eLCS improvements relating to recall were very small as changes were made to the baseline but the results for positive in heart disease are very low from 0% to 0.06% for preprocessing and from 0.06% to 0.93% for the improved eLCS. Although these eLCS configurations almost predict all instances without heart disease it lacked the ability to detect a considerable amount of positive cases.

While the other model compared to the improved eLCS model showed a greater percentage of true positives. Logistic regression showed a massive jump of true positive detection from 0.93% to 28.44% when compared with the improved eLCS model while decision tree showed a 53.58% and random forest with a true positive rate of 64.45%. These values relate to Figure 6.3.1 of the recall summary metric. The rate of false positives increases to 3.35% (logistic regression), 6.07% (decision tree) and 6.88% (random forest).

These models clearly showed that the eLCS is not the best in terms of recall as the improvements made only slightly increased the positive detection rate while the chosen 3 conventional models have a considerably higher detection rate, though with more false positives. The evaluation and explanation would be discussed in task 7.




Figure 6.3.3: Training time comparison graph

Stated by Figure 6.3.3, the improved eLCS model had the longest mean training time of 224.39 seconds. The original baseline eLCS and preprocessed eLCS took 13.09 seconds and 11.56 seconds respectively. All conventional models were trained in less than 2 seconds.

The improved eLCS trained much slower due to the hyperparameter tuning, which involved increasing the rule population to 8,000 and learning iterations to 50,000. These changes allowed more rules to be stored, compared and updated, while the iterations gave the system more opportunities to evaluate and refine those rules. Undersampling did reduce the records, but wasn't enough to affect the training time as much as the hyperparameters.

The improved eLCS model was around 17 times slower than the baseline eLCS model with minimal preprocessing, and around 19 times slower than the original eLCS model with full preprocessing. Comparing these three metrics, it can be stated that parameter tuning had the greatest influence on training time for each experiment.

The main comparison we care about is the improved eLCS model with the conventional models, because they use the same preprocessed and undersampled training sets. The difference in time between these models was significant with the improved eLCS model being 909 times slower than logistic regression, 684 times slower than decision tree, and 127 times slower than random forest. Out of the three conventional models, random forest was the slowest. However, out of all 6 experiments it achieved the highest recall, F1 and balanced accuracy results (Figure 6.3.1). It also achieved the same ROC-AUC as the improved eLCS model, while requiring significantly less time.




Figure 6.3.4: Mean precision recall curve across 10 seeds

This precision-recall curve shows the mean precision across 10 seeds, with the average precision of each model in the graph. All the models stayed above the test set mean prevalence of 2.11% for most recall values until 1. However, the mean precision is still considered low, with the range being from 0.038 to 0.183 across all models.

The beginning recall values show varying amounts of precision decrease across all the models. There is a significant decrease in precision for most models at the beginning, from 0.0 to 0.2. The baseline eLCS model shows the slightest steep decrease in precision; however, it has the highest deviation in precision of 0.046. Other models, such as random forest and improved eLCS precision dropping almost instantaneously. The random forest and improved eLCS precision dropped from 1 to 0.25 and level on about 0.2 for most of the recall points. Preprocessed has the lowest mean precision across the board at 0.037. The logistic regression is quite the opposite compared to the rest of the models, with precision increasing until 0.2, after which it starts to decline.

From 0.4 to 1.0 recall, all models show a slight gradual decrease until they reach the convergence point with the test set mean prevalence. However, the models vary in maintaining their precision throughout the recall points. The random forest showed the highest average precision of 0.167753 when considering the deviation, as well as other models that have a higher average precision but have a much higher deviation of precision. Through this comparison with the other models, the improved eLCS showed a more consistent mean precision than other models and other configurations of eLCS, such as the baseline model, original preprocessed model, logistic regression, and decision tree, but random forest remains the most competitive model, as it maintained a high mean precision with low variation; no model has the highest precision at every recall percentage.

7. Interpretation of Results and Explainable AI

The conventional models suggested in section 6 were much more successful at detecting positive heart disease cases than the eLCS models. Random forest had the best overall performance with a balanced accuracy of 0.788 and a recall of 0.644. The logistic regression and decision tree also showed much higher recall as compared to our improved eLCS models. On the other hand, the improved eLCS remained competitive with ROC-AUC and PR-AUC. This means that the eLCS was better at finding relationships and patterns between the two classes, however it struggled to convert these into positive predictions.

This shows that the main value from eLCS for this project is not that it achieved the best performance for prediction, but instead being able to have explainable patterns, learnt by its rule based structure.

7.1 Inspection of the learned rule population

The improved eLCS model was inspected after training to see the classifiers stored in the population. The analysis only focused on classifiers who predicted the patients having a heart disease, which is target class 1.

For the specific rules being inspected, instead of selecting them purely based on high accuracy, the rules were sorted by their match count. This gives rules that have been used across larger training instances more importance. This is because some rules had an accuracy of 1.0, however they were matched in only two records, this provides very little information on how credible the rule is.




Figure 7.11: Inspection of the final improved eLCS rule population



Figure 7.12: Internal values of an example class-1 eLCS classifier

7.2 Interpretation of learned heart disease rules

In eLCS, each classifier contains a genotype and phenotype which is a list of conditions and its predicted class. When a patient matches with the classifier's conditions, then the classifier may contribute to the final prediction. Many classifiers may contribute to the same prediction, so an individual rule should not be considered as the main reasoning for the model's prediction.

The top three genotype / classifier rules ordered by their match count were selected for further inspection. All three of them predicted positive for heart disease, however they all have different features.

Rule	Attributes	Conditions	Prediction	Accuracy	Fitness	Match Count	Correct Count	Numerosity
248 650	alcohol_intake_binned, smoking_Current, high_b...	[1.0, 0.0, 1.0, 3.0, 1.0, 2.0, 1.0, 2.0, 1.0, ...	1	0.622708	0.093632	1418	883	1
271 699	systolic_bp_binned, age_binned, gender, alcoho...	[2.0, 3.0, 1.0, 1.0, 0.0, 1.0, 1.0, 1.0]	1	0.627702	0.097447	1249	784	1
242 637	systolic_bp_binned, gender, alcohol_intake_bin...	[2.0, 1.0, 1.0, 2.0, 3.0, 1.0]	1	0.627641	0.097399	1136	713	1
363 888	smoking_Current, age_binned, ldl_binned, diast...	[0.0, 3.0, 1.0, 1.0, 0.0, 1.0, 1.0, 2.0]	1	0.634015	0.102447	929	589	1
634 1497	systolic_bp_binned, crp_level_binned, choleste...	[2.0, 2.0, 1.0, 0.0, 2.0, 1.0, 3.0, 0.0, 1.0]	1	0.639522	0.106974	921	589	1
301 765	employment_status_Unemployed, gender, heart_ra...	[0.0, 1.0, 1.0, 0.0, 1.0, 0.0, 1.0, 2.0, 0.0, ...	1	0.675894	0.141056	867	586	1
632 1494	age_binned, systolic_bp_binned, employment_sta...	[3.0, 2.0, 0.0, 1.0, 1.0, 1.0, 1.0, 0.0, 1.0, ...	1	0.627240	0.097089	837	525	1
176 507	smoking_Former, diabetes, diastolic_bp_binned,...	[0.0, 1.0, 1.0, 3.0, 2.0, 2.0, 0.0]	1	0.645161	0.111774	775	500	1
330 821	hdl_binned, crp_level_binned, gender, diabetes...	[2.0, 2.0, 1.0, 1.0, 0.0, 1.0, 2.0, 0.0, 1.0]	1	0.631158	0.100159	751	474	1
315 794	systolic_bp_binned, diastolic_bp_binned, age_b...	[2.0, 1.0, 3.0, 1.0, 2.0, 1.0, 0.0, 0.0]	1	0.658177	0.123513	746	491	1

Table 7.2.1: Examples of high-experience heart-disease rules learned by eLCS

Rule 650 had the most matches, matching 1,418 training instances and correctly classified 883 of them having a heart disease. The rule is a combination of age, blood pressure, cholesterol, inflammation, salt and alcohol consumption. From inspecting the genotype of classifiers more closely, we can see how eLCS is interpretable to humans, and how many different features combine to form rules for predictions.

Rule 699 matched 1,249 instances and was correct for 784 of them. This rule is slightly less complicated and includes senior age, male gender, Stage 1 blood pressure, low alcohol and salt intake, and normal heart rate.

Lastly, **rule 637** matched 1136 instances and made 713 correct classifications. It was the easiest rule of the three, and composed of senior age, male gender, overweight BMI, Stage 1 systolic blood pressure, and medium alcohol and salt intake.

Comparing these rules together, we can see there are some similar combinations amongst them. There are also some similar combinations in the rules. All the three included senior age, Stage 1 diastolic blood pressure, medium alcohol intake and medium salt intake.

However, these do not be interpreted as evidence that these features are associated with heart disease. They are relationships the eLCS has learnt after 50,000 iterations based on the dataset. They are the combinations the model associates with positive cases.

7.3 Interpretability and trustworthiness

The main benefits of eLCS is that it is not a black box model, but an interpretable model. All the learned knowledge by the model can be inspected as a set of understandable if else rules. It is straightforward for humans to understand the reasoning behind the model's predictions and decisions. To see which features are being used together in each rule, and which class each rule supports. This is often not the case for models such as random forest where its underlying reasoning is difficult for humans to understand.

However, these rules do not guarantee that the model is reliable. The rules that were inspected above were only around 62-63% accurate when they matched. And from section 6, the improved eLCS model had an extremely low overall recall when compared to the conventional models. Therefore, the rules provide a good explanation for what the model has learnt, but they do not necessarily produce reliable enough predictions for real world clinical use.


Additionally, the rules were learned from this specific data set, which may not generalise to other patients. This means that they should be treated as patterns made by the model rather than medical advice. Testing and reviewing from medical experts would be required before any of these rules could be considered medically meaningful.

8. Discussion, Limitations, and Future Work

The goal of this project was to investigate whether preprocessing, class balancing, and hyperparameter tuning could improve an eLCS model and predicting heart disease. The actions used produced mixed results where some aspects improved while others went down. After applying sampling, the model's positive-case detection improved immediately. Hyperparameter tuning improved the ranking metrics but also reduced the recall and balance accuracy. In the end, random forest achieved the strongest overall classification performance with the highest F1, recall, and balanced accuracy scores. However, none of the models are ready to be used in real-life detection scenarios.


8.1 Reasons for performance improvement (or lack thereof)

The original dataset performed poorly because it contained class imbalance., where 97.9% of the records belonged to class 0, and 2.1% to class 1. Because the baseline eLCS model mainly observed class 0 records, it learned to pick the majority class every time. This is why the recall, precision and f1 were zero for the baseline model. The 97.9% accuracy was misleading, because if the model predicts everything as negative for heart disease, it achieves the same accuracy. Because of these findings, the baseline model could not detect heart disease.



Preprocessing removed all unnecessary data including the missing values, duplicate rows, and misleading attributes. Attributes like sublabel, disease_flags, heart_disease, and label were removed due to leaking the answer off the target attribute. Categorical variables were encoded consistently, and numerical variables were converted into categorical groups in order to minimise the number of conditions. All of these changes allowed the data to be more interpretable by the eLCS model, however, it did not solve the class imbalance issue. Recall increased by only went from 0 to 0.0006, and the balanced accuracy went from 0.500 to 0.5003. This indicates that the preprocessing applied did not solve the class imbalance, because class 1 was still approximately 2.1% of the dataset.

To solve class imbalance, sampling was used to fairly reduce the class 0 population. After randomly picking 25,000 class 0 records, the ratio of negative and positive records improved from 46:1 to approximately 5:1. This allowed positive cases to have more influence during the rule learning process. During experimentation, the eLCS model detected positive cases more frequently and recall



increased from approximately 0.0006 to 0.0502. The balanced accuracy and PR-AUC also increased from 0.5003 to 0.5229 and 0.0376 to 0.1016 respectively. Overall undersampling didn't solve all problems, as class imbalance still existed with a ratio of 5.3 negative cases for every 1 positive case, causing the model to observe more class 0 records.

Hyperparameter tuning was used to change the model's rule population and learning iterations. The rule population was increased from 1,000 to 8,000, while the learning iterations was increased to 50,000 due to the undersampled training set containing approximately 30,000 records. The increased rule population allowed more rules to be stored, while more iterations allowed rules to be evaluated and refined more frequently. The result of this was a more conservative model where the ranking improved, but less rankings were stated as positive cases. This means that the selected values for population and learning iterations were better for ranking each record, but worse for predicting positive cases as the final decision.

Random Forest performed the highest out of all models, with a recall of 0.644, balanced accuracy of 0.788, F1 score of 0.267, and a ROC-AUC of 0.927. These were all higher than the improved eLCS model scores for the same variables, except for the ROC-AUC which was the same. It also finished training in 1.77 seconds compared





Figure 8.1.1 : Confusion matrices of all models

to the 224.39 seconds it took for the improved eLCS model, making it more computationally efficient.

8.2 Practical significance of results beyond raw numbers

The original and preprocessed eLCS models were not useful for predicting heart disease cases at all. The high accuracy of 97.9% was misleading, because predicting every record as class 0 resulted in that accuracy. The improved model after parameter tuning and undersampling only managed to predict 0.93% of positive cases, meaning that more than 99% of the positive cases went undetected. In healthcare, missing diseases can be life-threatening. Therefore, the improved eLCS model is not acceptable even with a high ROC-AUC value.

The ROC-AUC value of 0.926 indicates that the model can rank positive cases above negative cases, and that there were relationships learnt during training. However, the ranking alone doesn't mean the final predictions are useful, and the prediction behaviour of the model was too conservative. This finding is important because usually a high ROC-AUC makes the system appear as successful.




The improved eLCS model had a precision of approximately 0.260 which means that around 26% of the model's predictions were correct. This was the highest of any model trained, meaning that the improved eLCS model had the highest true/false positive ratio. However, the improved eLCS model also made very little predictions. The high precision is good but it doesn't compensate for missing over 99% of positive cases. False negatives are a **type II error** in the medical field, which is the **most dangerous error** when it comes to predicting a patient's status. Precision needs to be interpreted with recall and the confusion matrix for us to see if the model performs well.

Random Forest identified approximately 64.45% of all positive cases. This was considerably better than the improved eLCS model, but it still missed 35.55% of actual positive cases. The cost of the improved recall



value was that more false positive predictions were made. In the context of real life disease screening, false positives are more acceptable than false negatives because a life-threatening disease is not going unnoticed. Furthermore, the patient won't actually have the disease in a false positive. This is a **type I error** because the prediction is wrong, and could cause major stress on the patient and their close ones, increased healthcare costs, and more clinical workload for the hospital.

All of the models trained in this project are research prototypes, and should not be used to give a final decision for a heart-disease diagnosis. Even the most efficient models should not be fully relied upon for these critical decisions, but as a tool to support a professional. As new models arise and get more efficient at detecting heart-disease or any type of diseases, the health professional is always responsible for the final decision.




8.3 Limitations

Class imbalance was a major limitation included in the dataset. Only around 2.1% of the dataset records represented heart disease cases, while 97.9% represented negative cases. On top of this, the dataset contained over 280,000 patient records. After splitting the dataset into training and testing splits, the training set contained only 4.738 positive cases. Class imbalance was the main problem during model training as it made accuracy unreliable, and limited the number of class 1 patterns available to the eLCS model.


Undersampling was the second limitation we faced. Although sampling was used in the improved eLCS model, it never fully went away. Taking away so many real records could have removed negative records that contained useful patterns for the eLCS model. Another problem with sampling is that different selections can create different rule populations which can affect how the model behaves during the decision making process. Only one sample was used for testing, because different samples could produce different results, we cannot say how effective it is to the model unless more samples are used for training.

The dataset was made up of three different source datasets, relating to heart-disease, hypertension, and diabetes. These sources may have used different data collection techniques and conditions without us knowing. While removing the source_dataset prevented shortcut learning, there could be differences between the source populations in other attributes. All clinical ranges were dealt with, however we cannot say whether the included values were accurately measured as data entry errors could still exist.



Missing values wasn't an issue during preprocessing, because the dataset didn't contain any. However, this shouldn't be interpreted as the dataset having perfect data quality. The sources could have removed, filled or transformed any missing values that existed, which could have been performed with bias. Duplicate records were removed in the first assignment. However, task 3 found repeated combinations with some of them having conflicting target values. These were kept because removing them could have removed valid minority class data.

During the last assignment we found that individually, features had a weak relationship with heart disease. This made it difficult for eLCS to find rules that were reliable. Using feature engineering, the data was simplified to improve rule readability, but some actions like discretising numerical features could have removed useful information. Ordinal encoding was also used, but it assumes that categorical attributes have an ordered relationship with heart disease. Therefore some feature-engineering choices could have limited the model's performance.



Hyperparameter tuning was applied with only one iteration, with rule population = 8,000 and learning iterations = 50,000. These settings decreased the model's recall and balanced accuracy scores, from undersampling. This doesn't mean that changing the hyperparameters are bad, but that more settings could be explored. While reducing the recall and balanced accuracy, it did increase the precision and ROC-AUC scores.

8.4 Future work

As found during the comparison of the different models with the improved LCS developed in task 4, the improved LCS model consisting of undersampling and hyperparameter changes did not substantially improve positive case detection when compared to the baseline and preprocessed eLCS. Compared to the other selected conventional models, it did not perform as well as the others, as the conventional models achieved a higher recall, balanced accuracy, and F1 score.

In the current improved system, it was correct to identify that class imbalance was a massive issue in our dataset, in which we used undersampling at an approx. 5.3:1 ratio. However, it is still heavily favoured towards the majority class. Therefore, a more aggressive undersampling method could be used, such as a 2:1 ratio or 1:1 ratio, or utilize oversampling.

The hyperparameters could be experimented with by changing the crossover probability, mutation probability, learning iterations, and population size. It is important to systematically evaluate each change so that whether it adds any meaningful improvements can be identified. This would also affect the time it takes for training, as the current improved eLCS model takes an incredibly long training time compared to other models; this is likely due to the 50,000 training iterations and 8000-rule population. Changing these could find a better balance between positive case detection and computational cost.

The preprocessing phase could be further reviewed and made sure that no meaningful information is removed due to binning. So potentially retaining continuous values and applying scaling could be more beneficial than the current implementation of discretization of numerical values and the encoding of both the numerical and categorical values. This would help determine if preserving numerical data improves the eLCS performance.

Lastly, individuals with domain knowledge about heart disease, such as a cardiologist, could validate the source and help validate whether the dataset is accurate. And check that the most-matched rules found, such as rule 650, are sensible, and whether the classifier found is usable in the real world.

Section 9: Contribution and Reflection

9.1 Contribution Statement

Name	Description	Total Contribution %
Fateh	• Baseline eLCS on raw dataset • Feature engineering • Improved LCS-based system • Model comparison • Discussion and Limitations	33%
Leo	• Baseline eLCS on raw dataset • Data Preprocessing and Feature Engineering • Experiment design and evaluation metrics • Statistical significance testing • Interpretation of rules and Explainable AI	33%
Shawn	• Preprocessing • Intro and justification of eLCS • Comparison with 3 other conventional models • Future work needed	33%

9.2 Individual Reflections

Fateh (23217987): Working on this assignment improved my understanding of data engineering and the process of machine learning. I learned how to perform data preprocessing, evaluate models, sampling, hyperparameter tuning and analysing experimental metrics. I learned that just because model scores are high, doesn't mean that the model can be used in the real world. In fact, all variables must be interpreted together before stating how the model will perform with real cases. I believe that I developed stronger skills relating to machine learning, and am confident in my ability to train models.

Shawn (23204035): This assignment has helped me understand the unsupervised / reinforced learning eLCS model by actively implementing it into a real world problem. Throughout the process, I learned how to understand the dataset and try to improve it during the preprocessing stage, especially by identifying data leakage, class imbalance, and transformation of the data. As well as trying to improve the eLCS from the baseline eLCS. I also learned during the comparison that not all novel changes would improve our model, as shown in task 6 of the comparison.

Leo (23277188): This assignment has improved my understanding of the entire machine learning process, from preprocessing, effects of tuning hyperparameters, to evaluating and comparing different models. It has also taught me how different factors like population size, number of learning iterations, and class balancing influence the eLCS algorithm and how to properly evaluate models with unbalanced data. I also learnt how eLCS's rule based structure makes its predictions much easier to interpret than other algorithms, but we should still be cautious in trusting the predictions made.

References

Dataset

- <https://www.kaggle.com/datasets/danishjmeo/281k-patient-dataset-for-multi-disease-prediction?resource=download>

Sub Datasets

- <https://www.kaggle.com/datasets/iammustafatz/diabetes-prediction-dataset>,
- [kaggle.com/datasets/oktayrdeki/heart-disease](https://www.kaggle.com/datasets/oktayrdeki/heart-disease),
- [kaggle.com/datasets/ankushpanday1/hypertension-risk-prediction-dataset](https://www.kaggle.com/datasets/ankushpanday1/hypertension-risk-prediction-dataset)

Research

Latest heart disease statistics. (n.d.). Heart Foundation NZ.
<https://www.heartfoundation.org.nz/statistics>

GitHub Link

<https://github.com/fatehbhular/Data-Engineering-and-AI-Project.git>