

Data Engineering and Machine Learning Pipeline

Phase II Project Report

ENGE707 Data Engineering and AI

Leo Cao (23211788), Fateh Bhular (23217987), Shawn Lee (23204035)

Auckland University of Technology

Table of Content

1. Introduction – Project Phase-I Summary.....	4
1.1 Problem Restatement.....	5
1.2 Dataset summary.....	5
1.3 Data-quality issues identified in Phase I.....	6
1.4 Justification for an LCS-based solution.....	6
2. Original LCS System on Raw Dataset.....	7
2.1 LCS variant selection (preferably eLCS) – justification.....	7
2.2 Confirmation that original LCS code is used without algorithmic modification.....	8
2.3 Execution on raw / minimally processed dataset.....	9
2.4 LCS parameter settings used.....	11
2.5 Baseline performance results.....	12
2.6 Minimal processing required for LCS compatibility.....	13
3. Data Preprocessing and Feature Engineering.....	14
3.1 Data-leakage prevention measures.....	14
3.2 Data cleaning.....	15
3.3 Outlier investigation and handling.....	18
3.4 Feature engineering.....	19
3.5 Encoding categorical variables.....	20
3.6 Discretising & Encoding numerical variables.....	20
3.7 Class-imbalance handling.....	21
3.8 Expected impact of preprocessing on LCS performance.....	22
4. Improved LCS-Based System.....	22
4.1 Summary of changes vs. the original LCS system.....	23
4.2 Justification – expected effect on accuracy or model quality.....	23
4.3 Algorithmic explanation (pseudocode or flowchart).....	23
4.4 Comparison against the original LCS baseline.....	23
5. Experiments.....	23
5.1 Validation strategy (80/20 split or k-fold CV) – justification.....	24

5.2 Evaluation metrics – selection and justification (classification: accuracy, precision, recall, F1, balanced accuracy, ROC-AUC, PR-AUC, confusion matrix; regression: MAE, MSE, RMSE, MAPE, R ² , residual analysis).....	24
5.3 Statistical significance testing.....	24
5.4 Interpretation of significance results.....	24
6. Model Comparison.....	24
6.1 Original LCS – raw/minimally processed data.....	25
6.2 Original LCS – preprocessed data.....	25
6.3 Improved LCS – preprocessed/feature-engineered data.....	25
6.4 At least three conventional ML models (e.g. Random Forest, SVM, Decision Tree, Logistic Regression, Naive Bayes, k-NN, Gaussian Process) – no deep learning.....	25
6.5 Comparative results table.....	25
7. Interpretation of Results and Explainable AI.....	25
7.1 Interpretation of main findings from the model comparison.....	26
7.2 Explanation of key LCS rules / feature conditions / classifier patterns.....	26
7.3 At least three sample LCS rules discussed for decision-making insight.....	26
7.4 How rule-based outputs support interpretability.....	26
7.5 Trustworthiness and practical meaningfulness of the explanations.....	26
8. Discussion, Limitations, and Future Work.....	26
8.1 Reasons for performance improvement (or lack thereof).....	26
8.2 Practical significance of results beyond raw numbers.....	26
8.3 Limitations – data quality, dataset size, class imbalance, noise, missingness, feature relevance...26	
8.4 Future work – improved LCS variants, better preprocessing, external validation, domain expert review.....	26
References.....	27

1. Introduction – Project Phase-I Summary

- ☐ Students must briefly summarise the Phase I project. Students are required to:
 - ☒ ~~Restate the selected problem, dataset, target variable, and prediction objective.~~
 - ☒ ~~Summarise the dataset source, number of records, number of features, and main data characteristics.~~
 - ☐ Briefly describe the main data-quality issues identified in Phase I.
 - ☐ Explain why the problem is suitable for an LCS-based solution.

Link with justification, do not repeat

- raw data, no preprocessing
- base elcs model
- apply preprocessing, we want to see if our preprocessing was good or not
 - Evidence between pre processing and the processed data
 - reducing the search space -
- Make some changes in the ELCS base system to improve the accuracy or computational cost
e.g. lowering the training amount (Novelty)
 - changing fitness function is novelty
- Training/Testing sets should be run at least 10 times. take the average of 10 times
 - use different seed
 - *also mention the seeds in the report.*
- cross fold validation not random and it is in sequence
- **don't forget to randomize when separating testing and training data**

1.1 Problem Restatement

Heart disease is a major concern that currently affects 195,000 people in New Zealand alone. Approximately every 90 minutes, a New Zealander dies from heart disease, and it is the number one killer for women (*Latest Heart Disease Statistics*, n.d.). In this project, the dataset “281K Patient Dataset for Multi-Disease Prediction” from Kaggle is used. The goal is to find whether the lifestyle, demographics, and clinical measurements in the dataset have any relevance to predicting whether an individual has heart disease. A model with this prediction ability isn’t a full replacement for having an in-person cardiovascular risk assessment. Still, it could be a cost-effective way to check for an individual’s chance of having heart disease. The target column “heart_disease_target” is a binary classification created after cleaning and is based on the “sublabel” column, which contains the combination of diseases an instance has. If the instance from the sublabel contains “HT”, the value of the new column would be True; otherwise, False. This new column is created to be easily used to compare and train the machine learning model’s prediction capabilities.

1.2 Dataset summary

The dataset is a combination of three datasets (diabetes, heart disease, and hypertension). There are 280,985 raw records, and 278,701 remain after removing duplicates; the 278,701 will be used instead of the raw amount. The original dataset has 39 columns, but the introduction of “heart_disease_target” gives 40 columns. The columns utilized during the analysis and processing stage are lifestyle, demographics, clinical, and disease indicators. No missing values were found during the dataset inspection.

The overall heart disease rate is 2.1% across all these sources, but individually the dataset shows varying heart disease rates: 20.06% in the heart disease dataset source, compared to 4.17% for the diabetes source, and 0% for the hypertension source.

1.3 Data-quality issues identified in Phase I

From the data inspection and data cleaning stage, multiple issues in the dataset were identified:

- 2,284 duplicate instances were removed.
- heart_disease has fractional values (from 0 to 1), which doesn't clearly explain whether an individual has heart disease, so a new column (heart_disease_target) was created to serve as the target (binary) instead of the heart_disease column.
- The columns sublabel, disease_flags, label, and heart_disease all contain data leakage information. This provides clues and partial answers to whether the individual has heart disease; thus, they are excluded from the prediction process.
- Some columns such as age/age_level/age_normalized and bmi/bmi_level carry the same information; this is to be reviewed before modeling.
- Imbalance in the distribution of instances with and without heart disease, since 97.9% of records do not have heart disease and 2.1% do.
- As described in the dataset summary, individual sources have different heart disease rates.

1.4 Justification for an LCS-based solution

From the exploratory data analysis in phase 1, the heat map demonstrated that there was a weak linear correlation with the heart_disease_target; e.g age is 0.08, BMI is 0.04, and triglycerides are 0.09. This suggests that a single feature alone does not have a strong correlation with the target. When just analysing BMI level and age group, a clear difference in heart disease rate was observed between different age groups and BMI levels; seniors who are also overweight displayed the highest heart disease rate of 6.53% compared to 2.1% of the overall heart disease rate. Adding gender results in a difference in heart disease rates of 8.18% for males and 4.99% for females who are seniors and overweight. The overall takeaway is that single features are weak as linear predictors, and that combining specific features together showed a greater heart disease rate.

The main reason why LCS is suitable for this problem is that it allows us to find a specific set of conditions that predicts the heart_disease_target well. In the exploratory data analysis, all possible combinations of features with the target could not be examined due to time constraints as well as the number of possible combinations of features in the dataset; therefore, LCS would help with automation. Another reason is that LCS involves a population to create if-then rules (classifiers). This could allow interpretability for each classifier, as the produced model aims to help predict heart disease. Experts in the field could analyse these attributes to see whether each classifier makes medical sense. However, there is a challenge: the dataset has a clear imbalance of instances with heart disease compared to those without; thus, the LCS machine learning model may favour the majority class. Further evaluation would be needed beyond just the accuracy.

2. Original LCS System on Raw Dataset

- ☐ Students must run an original LCS system on the raw or original dataset before making improvements. Students are required to:
 - ☐ Select an appropriate LCS variant, preferably eLCS.
 - ☐ Use the original LCS code without algorithmic modification.
 - ☐ Run the original LCS system on the raw or minimally processed dataset.
 - ☐ Report the selected LCS parameters.
 - ☐ Record the baseline performance results.
 - ☐ Explain any minimal processing required to make the dataset compatible with the LCS code.

2.1 LCS variant selection (preferably eLCS) – justification

We used eLCS, also known as Education Learning Classifier System, as our LCS system for this project. eLCS is a variant of Learning Classifier systems that uses supervised learning, and is used for classification problems. This project's goal is to successfully predict whether a patient has heart disease, which is represented by "heart_disease_target" with a value of True or False. A True/False value can be represented by binary 1 for True and 0 for False, making it a binary classification problem.

In our dataset “281K patient dataset”, there are both continuous and categorical variables. eLCS can work with both of these types of variables which made it a suitable LCS for this project.

We learned in phase I that there were no relationships between individual features and the target attribute. However, after combining features together and then comparing them, the relationships were clearer. eLCS is capable of learning rules that include a combination of attributes, which is exactly what this project requires for a machine learning model.

2.2 Confirmation that original LCS code is used without algorithmic modification

```
setup(  
    name = 'scikit-eLCS',  
    packages = ['skeLCS'],  
    version = '1.2.4',  
    license='License :: OSI Approved :: GNU General Public License v3 (GPLv3)',  
    description = 'Educational Learning Classifier System',
```

Figure 2.2.1: eLCS version confirmation

Our original dataset was given to the original eLCS implementation. The version used was 1.2.4, and there were no modifications made to the source code or behaviour to the system. The setup remained untouched and was found in the “setup.py” file in the system files. External code was added outside the eLCS implementation package, but this was just to make the dataset ready for model training. The following code was added:

1. Loading the dataset
2. Creating the target attribute
3. Splitting the dataset into training and testing sets (80/20 split)
4. Predicting the results
5. Calculating all of the metrics from model training

2.3 Execution on raw / minimally processed dataset

The raw dataset from “patient_data.csv” was used for the baseline eLCS model. As shown in Figure 2.3.1, there were 280,985 records and 39 rows. The raw dataset contained 2284 duplicate records, which were intentionally kept. The duplicate rows will be handled during Task 3 because the original eLCS model was meant to represent the raw dataset with none to very minimal preprocessing.

```
Raw shape: (280985, 39)
Raw records: 280985
Raw columns: 39
Missing values: 0
Duplicate records: 2284
```

Figure 2.3.1:

Raw Dataset

heart_disease_target			
0	275081		
1	5904		
Name: count, dtype: int64			
heart_disease_target			
	0	1	
sublabel			
DI	31918	0	
DI_HT	0	1398	
DI_HT_HY	0	845	
DI_HY	66473	0	
HT	0	2595	
HT_HY	0	1066	
HY	69806	0	
N	106884	0	

The target attribute was not in the original dataset, so it was added after loading. The “sublabel” attribute was used to create it by checking if the value contained “HT”, which represented heart disease. The values including “HT” were mapped to class 1, while the values not including it were mapped to class 0.

The result of the target creation was that 275,081 records were mapped to class 0 (no heart disease), and 5,904 records were mapped to class 1 (heart disease). With these results, the positive class rate was approximately 2.10%.

Figure 2.3.2: Target Construction

After the target attribute was added, there were a total of 40 columns. However, 4 of these were excluded from the predictor matrix due to having a one-to-one correlation

```
Predictor shape: (280985, 36)
Target shape: (280985,)
Positive-class rate: 0.021011797782799792
```

Figure 2.3.3: Predictor Set

with the target. Therefore, the baseline experiment used 36 columns. We performed validation checks which confirmed the target attribute was binary and the dataset contained no missing values.

The dataset was split into training and testing sets with an 80/20 split. Stratification was used to preserve the class imbalance in both sets so that they could both be as close to the original dataset split as possible.

```
Training shape: (224788, 36)
Testing shape: (56197, 36)
Training positive rate: 0.021010908055590156
Testing positive rate: 0.021015356691638343
```

Figure 2.3.4: 80/20 Split results

The results shown in Figure 2.3.4 tell us that the training set contained 224,788 rows (80% of rows from the dataset) and the testing set contained 56197 rows (20% of the dataset). Both sets used 36 attributes because 4 columns from the original dataset had a one-to-one correlation with the target attribute. Both sets maintained the same positive rate of approximately 2.10%, which was the same rate in the raw dataset.

```
Starting seed 3..
Completed seed 3: balanced accuracy=0.5000, ROC-AUC=0.6925, PR-AUC=0.2241, positive predictions=0
Starting seed 6..
Completed seed 6: balanced accuracy=0.5000, ROC-AUC=0.7354, PR-AUC=0.1711, positive predictions=0
Starting seed 9..
Completed seed 9: balanced accuracy=0.5000, ROC-AUC=0.6086, PR-AUC=0.0782, positive predictions=0
Starting seed 12..
Completed seed 12: balanced accuracy=0.5000, ROC-AUC=0.6885, PR-AUC=0.1787, positive predictions=0
Starting seed 15..
Completed seed 15: balanced accuracy=0.5000, ROC-AUC=0.7252, PR-AUC=0.2188, positive predictions=0
Starting seed 18..
Completed seed 18: balanced accuracy=0.5000, ROC-AUC=0.6970, PR-AUC=0.2439, positive predictions=0
Starting seed 21..
Completed seed 21: balanced accuracy=0.5000, ROC-AUC=0.6830, PR-AUC=0.1565, positive predictions=0
Starting seed 24..
Completed seed 24: balanced accuracy=0.5000, ROC-AUC=0.7186, PR-AUC=0.1668, positive predictions=0
Starting seed 27..
Completed seed 27: balanced accuracy=0.5000, ROC-AUC=0.6563, PR-AUC=0.2024, positive predictions=0
Starting seed 30..
Completed seed 30: balanced accuracy=0.5000, ROC-AUC=0.6630, PR-AUC=0.1845, positive predictions=0
```

Figure 2.3.5: Results of the 10 baseline experiments

Figure 2.3.5 shows the results of the 10 repeated experiments using the original eLCS model. The seeds used were [3, 6, 9, 12, 15, 18, 21, 24, 27, 30], verifying that each experiment was unique. In every run, the seed affected both the shuffled split (pseudo-random number generator) and the random state of the eLCS. Besides the seed, all of the parameters were kept the same throughout all 10 experiments. We will use the same seeds in Part 4 for the improved eLCS model to make the comparison between data fair.

2.4 LCS parameter settings used

	Parameter	Value
0	N	1000
1	acc_sub	0.99
2	beta	0.2
3	chi	0.8
4	delta	0.1
5	discrete_attribute_limit	10
6	do_GA_subsumption	True
7	do_correct_set_subsumption	False
8	fitness_reduction	0.1
9	init_fit	0.01
10	learning_iterations	10000
11	match_for_missingness	False
12	mu	0.04
13	nu	5
14	p_spec	0.5
15	random_state	3
16	reboot_filename	None
17	selection_method	tournament
18	specified_attributes	[]
19	theta_GA	25
20	theta_del	20
21	theta_sel	0.5
22	theta_sub	20
23	track_accuracy_while_fit	False

Figure 2.4.1: Parameters List

When creating the baseline model, we used the seed as the random_state. So all **parameters were kept the same except for the seed** for each baseline experiment.

The parameters in Figure 2.4.1 were retrieved using “model.get_params()”, which listed all of the parameters from the newly created eLCS object. Each of these affects how the experiment runs; for example, “learning_iterations” represents how many times eLCS performs learning iterations, “N” represents the maximum rule-population size, and “mu” represents the crossover probability. The rest of the explanations are in “Original_eLCS_Baseline.ipynb” in the project files.

```
# 10 reproducible seeds for the repeated experiments.  
SEEDS = [3, 6, 9, 12, 15, 18, 21, 24, 27, 30]  
✓ 0.0s
```

Figure 2.4.2: Seeds list

We used 10 different seeds (Figure 2.4.2) in **10 different and unrelated baseline experiments**. For each seed, a new eLCS model object was created, where the seed affected the train-test set split and the randomness of the experiment. The seed used for each experiment was recorded in order to make each run reproducible. As shown in Figure 2.3.5, the seed was recorded along with the experiment results for each run.

2.5 Baseline performance results

	Metric	Value
0	Accuracy	0.978985
1	Precision	0.000000
2	Recall	0.000000
3	F1	0.000000
4	Balanced accuracy	0.500000
5	ROC-AUC	0.545988
6	PR-AUC	0.034085
7	Training seconds	12.129322

	Predicted 0	Predicted 1
Actual 0	55016	0
Actual 1	1181	0

Final training balanced accuracy: 0.5
Final training coverage: 1.0

	Seed	Accuracy	Precision	Recall	F1	Balanced accuracy	ROC-AUC	PR-AUC	True negatives	False positives	False negatives	True positives	Positive predictions	Training seconds
0	3	0.978985	0.0	0.0	0.0	0.5	0.545988	0.034085	55016	0	1181	0	0	12.364138
1	6	0.978985	0.0	0.0	0.0	0.5	0.606016	0.052822	55016	0	1181	0	0	12.287477
2	9	0.978985	0.0	0.0	0.0	0.5	0.553733	0.035440	55016	0	1181	0	0	12.408605
3	12	0.978985	0.0	0.0	0.0	0.5	0.570389	0.035298	55016	0	1181	0	0	12.864080
4	15	0.978985	0.0	0.0	0.0	0.5	0.603712	0.055804	55016	0	1181	0	0	12.837608
5	18	0.978985	0.0	0.0	0.0	0.5	0.649815	0.050285	55016	0	1181	0	0	12.465585
6	21	0.978985	0.0	0.0	0.0	0.5	0.522048	0.039263	55016	0	1181	0	0	12.424905
7	24	0.978985	0.0	0.0	0.0	0.5	0.568428	0.046623	55016	0	1181	0	0	12.537057
8	27	0.978985	0.0	0.0	0.0	0.5	0.613384	0.047341	55016	0	1181	0	0	12.737933
9	30	0.978985	0.0	0.0	0.0	0.5	0.576320	0.043518	55016	0	1181	0	0	12.252371

	Confusion category	mean	std	min	max
0	True negatives	55016.0	0.0	55016.0	55016.0
1	False positives	0.0	0.0	0.0	0.0
2	False negatives	1181.0	0.0	1181.0	1181.0
3	True positives	0.0	0.0	0.0	0.0

2.6 Minimal processing required for LCS compatibility

```
# all labels that include heart disease
HEART_DISEASE_LABELS = [
    "HT",
    "DI_HT",
    "HT_HY",
    "DI_HT_HY",
]

# copies the raw data and creates the binary target
df_baseline = df_raw.copy()

df_baseline["heart_disease_target"] = (
    df_baseline["sublabel"]
    .isin(HEART_DISEASE_LABELS)
    .astype(int)
)
```

```
# the target-related columns that cannot be predictors
EXCLUDED_COLUMNS = [
    "heart_disease_target",
    "sublabel",
    "disease_flags",
    "label",
    "heart_disease",
    "composite_key",
    "source_dataset",
]

# separates the predictor variables from the target
X = df_baseline.drop(columns=EXCLUDED_COLUMNS)
y = df_baseline["heart_disease_target"].astype(int)
```

```
# create a shuffled and stratified 80/20 train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.20,
    random_state=SEED,
    stratify=y,
    shuffle=True,
)
```

```
# learns category mappings from the training data and applies them to both partitions.
X_train_encoded[categorical_columns] = encoder.fit_transform(
    X_train[categorical_columns]
)
X_test_encoded[categorical_columns] = encoder.transform(
    X_test[categorical_columns]
)
```

3. Data Preprocessing and Feature Engineering

- ☐ Students must improve the quality of the dataset before developing the improved system and make it suitable for working with LCS. Students are required to:
- ☒ Address missing values, duplicate records, invalid values, inconsistent categories, and incorrect data types.
 - ☒ Investigate and handle outliers where appropriate.
 - ☐ Encode categorical variables and scale, discretise, or transform numerical variables where required.
 - ☐ Address class imbalance where relevant.
 - ☒ Apply feature engineering, feature selection, or dimensionality reduction where justified.
 - ☐ Clearly explain how preprocessing is expected to improve LCS performance.
 - ☒ Ensure that preprocessing avoids data leakage.

3.1 Data-leakage prevention measures

To prevent target leakage of heart_disease_target, there are multiple decisions made during the processing phase to address this issue. Two measures have been used to reduce the data-leakage, Train/test leakage and target leakage.

Train/Test leakage

```
Train shape: (222960, 39)
Test shape: (55741, 39)

Train target distribution:
heart_disease_target
False    0.978862
True     0.021138
Name: proportion, dtype: float64

Test target distribution:
heart_disease_target
False    0.978867
True     0.021133
Name: proportion, dtype: float64
```

The train/test split has been used before any other preprocessing measures. This is because splitting the data used for training and model testing should be divided. Models that are based on the same instances of data would cause the model to practically memorise the pattern instead of learning new patterns. A 80/20% split is used when 80%

of the data is used for training while 20% is used for testing(The 80/20 split would be further explained in **section 5**). Additional splitting the testing and training data stops the testing data from influencing the preprocessing phase, e.g calculating the median of glucose, as the testing data would change the median compared to just using the training data.

Target leakage

These columns label, sublabel, disease_flags, and heart_disease are removed in the processing stage because they have partial or full information about the target. The labels column classifies an individual as abnormal or normal with any type of disease, including heart disease, so this cannot be used. For the sublabel column, this was used to directly create the target because it contains the classification of diseases. Disease_flags column contains the binary indicator of which disease the individual has. Lastly, the heart_disease column shows the probability of heart disease for a patient with similar traits and showed a strong correlation of 0.65 with the target; therefore, it cannot be used.

```
columns_to_drop = ["sublabel", "disease_flags", "heart_disease", "label"] # All the columns to drop because it contains information about the target
X_train = X_train.drop(columns=columns_to_drop)
X_test = X_test.drop(columns=columns_to_drop)

print(X_train.columns.tolist())
✓ 0.0s
['composite_key', 'age_level', 'gender', 'bmi_level', 'smoking', 'diabetes', 'age', 'age_normalized', 'bmi', 'hypertension', 'HbA1c_level', 'glucose', 'c
```

(Note: **The rest of the processing is completed on the training dataset and then applied to the testing dataset**)

3.2 Data cleaning

In phase 1, the raw dataset has been cleaned for exploratory analysis (including missing values, duplicate records, inconsistent categories, and data types). Before going further, the dataset has been reviewed again to check if there are any issues after the splitting of the training and testing datasets. Invalid values were partially checked during phase 1 EDA; therefore, we will review them again.

Missing Values

```
Missing values in X_train: 0
Missing values in X_test: 0
```

Checked the missing values in the training set testing sets. No missing values were found. This is to ensure that the cleaning completed in phase 1 was

inspected again.

Duplicate Records

```
Duplicate rows in X_train: 165
Duplicate rows in X_test: 11
Target distribution among duplicate rows:
heart_disease_target
True      165
False     165
Name: count, dtype: int64

Example of conflicting duplicate pairs:

```

	composite_key	age_level	gender	bmi_level	smoking	diabetes	age
48110	Mature_Female_Overweight_Former_No	Mature	Female	Overweight	Former	No	51
26464	Mature_Female_Overweight_Former_No	Mature	Female	Overweight	Former	No	51
23489	Mature_Female_Overweight_Never_No	Mature	Female	Overweight	Never	No	47
28897	Mature_Female_Overweight_Never_No	Mature	Female	Overweight	Never	No	47

4 rows x 35 columns

When checking for duplicate records after taking measures for data leakage, removing the columns that have information about the target introduces 165 duplicate records in the training dataset and 11 records in the testing dataset. This is because the

target leakage columns contain values that make these instances unique. If the data were truly duplicated, then the target should not have conflicting results when the other columns have identical values. Therefore, these instances could not be removed since the different target values imply that these instances are different.

Invalid Values

	age	bmi	glucose	cholesterol
count	222960.000000	222960.000000	222960.000000	222960.000000
mean	49.712285	27.480704	135.341664	224.019387
std	21.692180	7.065210	38.508772	35.663319
min	2.000000	10.010000	70.000000	150.000000
25%	32.000000	22.000000	100.000000	201.000000
50%	50.000000	27.320000	138.000000	225.000000
75%	68.000000	32.700000	160.000000	243.000000
max	89.000000	95.220000	300.000000	300.000000

	systolic_bp	diastolic_bp	heart_rate
count	222960.000000	222960.000000	222960.000000
mean	134.444033	89.479164	74.509257
std	20.620661	13.735412	11.444668
min	90.000000	60.000000	50.000000
25%	125.000000	83.000000	69.000000
50%	134.167413	89.504587	74.487247
75%	144.000000	96.000000	80.000000
max	179.000000	119.000000	99.000000

Invalid values were checked by making sure that the values are within valid boundaries. Age, BMI, glucose, cholesterol, systolic_bp, diastolic_bp, and heart_rate all showed that the values are plausible.

Inconsistent categories

```
composite_key:
composite_key
Middle_Female_Overweight_Never_No    4502
Mature_Female_Overweight_Never_No    4355
Middle_Female_Obese_Never_No         4010
Mature_Female_Obese_Never_No         3990
Middle_Female_Normal_Never_No        3880
...
Young_Female_Underweight_Current_Yes  148
Young_Male_Underweight_Former_Yes     144
Young_Female_Underweight_Current_No   143
Young_Male_Underweight_Current_Yes    126
Young_Female_Underweight_Former_Yes   109
Name: count, Length: 192, dtype: int64

age_level:
age_level
Mature    69626
Middle    67504
Senior    55890
Young     29940
Name: count, dtype: int64
```

All categorical columns were checked for inconsistent categories. Composite key, age_level, gender, bmi_level, smoking, diabetes, physical_activity, family_history, stress_level, low_hdl_cholesterol, high_ldl_cholesterol, high_blood_pressure, sugar_consumption, education_level, employment_status and source_dataset, all have consistent categorical values and formatting consistency.


Data Types

```
composite_key    str
age_level        str
gender           str
bmi_level        str
smoking          str
diabetes         str
age             int64
age_normalized   float64
bmi              float64
hypertension     float64
HbA1c_level      float64
glucose          float64
cholesterol      float64
sleep_hours      float64
triglycerides    float64
physical_activity str
family_history   str
stress_level     str
low_hdl_cholesterol str
high_ldl_cholesterol str
blood_pressure   float64
high_blood_pressure str
sugar_consumption str
crp_level        float64
homocysteine_level float64
...
education_level  str
employment_status str
source_dataset   str
dtype: object
```

The data types of each column are correct; therefore, no changes would be made.

3.3 Outlier investigation and handling

```
age: lower=-22.00, upper=122.00, outliers=0
bmi: lower=5.95, upper=48.75, outliers=792
glucose: lower=10.00, upper=250.00, outliers=1626
cholesterol: lower=138.00, upper=306.00, outliers=0
triglycerides: lower=-42.62, upper=412.38, outliers=0
systolic_bp: lower=96.50, upper=172.50, outliers=21879
diastolic_bp: lower=63.50, upper=115.50, outliers=18624
heart_rate: lower=52.50, upper=96.50, outliers=16803
hdl: lower=34.50, upper=94.50, outliers=20099
ldl: lower=79.50, upper=179.50, outliers=23426
```



In phase 1, BMI was found to have a considerable number of outliers, but is proven that the outliers shouldn't be transformed or removed. The other columns were discussed at the surface level, based on the boxplot, and showed no isolated instances. To further investigate this, the Interquartile range method was performed on the numerical columns. Triglyceride, cholesterol, and age showed no outliers, but the rest of the columns' results displayed incredibly high values, e.g 18624 outliers for diastolic_bp as well as other columns. A histogram is created to show the distribution of instances across each column. Most of the columns (systolic_bp, diastolic_bp, heart_rate, hdl, ldl, glucose) showed a high peak in the middle and consistent values across the whole range. This is why using the IQR method and boxplot isn't enough to test for outliers and to show the outliers accurately. IQR is used only when the data shows a normal distribution; in this case, the data is not. It is concluded that these outliers would be treated as valid instances; thus, no removal, transformation, or capping is performed.

3.4 Feature engineering

In the current dataset up to this point of processing, multiple columns are irrelevant or provide no value towards predicting the target.

- Composite key **(removed)**

A composite key is just the combination of variables (age_level, gender, bmi_level, smoking, and diabetes); therefore, this does not provide any value, as it is simplifying group the variables for the instance. So removing this would not cause any information to be lost and is just an unnecessary column to provide to the model.

- Age_level and age_normalized **(removed)**

Age_level and age_normalized provide similar information to age, but age_level is a categorical level and age_normalized is the linear scaling of age and showed a 1.00 correlation with age during EDA. These add no additional information about age, as age could be discretised independently for eLCS, thus being a more favourable column to remain in the training dataset.

- bmi **(removed)**

The bmi column is removed because the bmi_level column follows the WHO standard, which is the standard classifier for BMI. The bmi column does allow independent interpretation to discretization, but since bmi_level is already in a suitable standard, there isn't any reason not to use bmi_level instead of bmi.

3.5 Encoding categorical variables

	31257	122755	162764	2022	249582
gender	0.0	1.0	0.0	0.0	0.0
diabetes	0.0	1.0	0.0	0.0	0.0
family_history	0.0	0.0	1.0	0.0	0.0
low_hdl_cholesterol	0.0	0.0	0.0	0.0	0.0
high_ldl_cholesterol	0.0	1.0	1.0	0.0	1.0
high_blood_pressure	0.0	1.0	0.0	0.0	1.0
bmi_level	3.0	2.0	3.0	0.0	3.0
smoking	0.0	0.0	0.0	0.0	1.0
physical_activity	0.0	1.0	1.0	2.0	1.0
stress_level	2.0	0.0	2.0	1.0	2.0
sugar_consumption	2.0	2.0	0.0	1.0	0.0

All the categorical variables have been encoded to numerical values because eLCS could not accept

values in strings. For example, gender has been changed from “Female” and “Male” to 0 and 1.

Multiple types of encoding have been used for the categorical variables since some variables required different encoding methods.

For these columns (gender, diabetes, family_history, low_hdl_cholesterol, high_ldl_cholesterol, and high_blood_pressure), binary encoding is used because these columns only have 2 possible outputs.

For these columns (bmi_level, smoking, physical_activity, stress_level, sugar_consumption, education_level), ordinal encoding is used because, for example, for stress_level, there is an order of the output values “low”, “medium,” and “high”; therefore, the encoding converts it into numerical values (1 to 3).

For the rest of the categorical columns (“employment_status”, “source_dataset”), one-hot encoding is used since they provide no order and cannot be represented in binary; therefore, new columns would be created to represent each possible outcome, as shown in the encoding categorical screenshot.

3.6 Discretising & Encoding numerical variables

	glucose_binned	cholesterol_binned	systolic_bp_binned	\	
31257	2.0	1.0	2.0		
122755	2.0	0.0	0.0		
162764	2.0	2.0	0.0		
2022	2.0	1.0	2.0		
249582	0.0	1.0	0.0		
	diastolic_bp_binned	HbA1c_level_binned	age_binned	\	
31257	1.0	1.0	2.0		
122755	0.0	2.0	1.0		
162764	0.0	0.0	3.0		
2022	1.0	0.0	1.0		
249582	0.0	0.0	1.0		
	triglycerides_binned	hdl_binned	ldl_binned	heart_rate_binned	\
31257	2.0	2.0	1.0	1.0	
122755	0.0	2.0	3.0	0.0	
162764	0.0	1.0	2.0	0.0	
2022	2.0	2.0	1.0	1.0	
249582	0.0	1.0	2.0	1.0	
	crp_level_binned	sleep_hours_binned	homocysteine_level_binned	\	
31257	2.0	1.0	0.0		
122755	2.0	1.0	0.0		
162764	2.0	0.0	0.0		
2022	2.0	0.0	0.0		
249582	2.0	1.0	0.0		
	hypertension_binned	alcohol_intake_binned	salt_intake_binned		
31257	1.0	1.0	1.0		

Discretising has been used for all the numerical datasets to turn them into categorical values because eLCS operates on if-then rules, which requires categorical data in order for the prediction to happen.

To determine the categorical values, the columns that have clinical standards are then used to determine whether they are low, medium, or high. For example, for glucose, the normal range would be below

or at 100mg/dL, the prediabetes level at 100 to 126mg/dL and diabetes at 126mg/dL and above. The columns (hypertension, alcohol_intake and salt_intake) that don't have a set standard of measuring levels require the use of quantile based discretisation. This is to roughly split all training data into 3 categorical values based on their relative numerical value from each other on each column. For example, hypertension in this dataset is a continuous range from 0-1. There is no standard measurement for this since hypertension is standardly measured in mmHg.

After discretising all the numerical data, these are then encoded so that eLCS could utilize them in the prediction.

3.7 Class-imbalance handling

In the EDA, the target variable is highly imbalanced across the dataset, with 2.1% of instances with heart disease and 97.9% of instances without heart disease of the total 278,701 instances. The baseline eLCS prediction results showed that the final balanced accuracy test is 0.5, this is effectively choosing negative (no heart disease) for all the instances. This result is no better than just predicting half of the instance does and half the instance does not when considering the balanced accuracy.

The imbalance would be handled after receiving the result from the prediction preprocessed dataset (**Task 4**) in eLCS. This is because it would be valuable insight into how much class imbalance could affect the overall prediction of the model when the imbalance is handled.

3.8 Expected impact of preprocessing on LCS performance

The overall preprocessing phase of the dataset is expected to improve the eLCS ability to process and performance, these improvements include:

- **Improved Computational cost / quicker prediction**

In the feature selection process, multiple redundant columns were removed, therefore there are less combinations of if and then rules.

- **More compact rule set**

Applying the discretization lowers the amount of classifiers since there will be less classifiers that are tied to continuous values which also increase readability for validation and inspection of the rule.

- **Improved accuracy of rules**

Using the appropriate encoding types would reduce the number of misleading information when just applying a general encoding type to all columns which causes the eLCS to learn inaccurate patterns.

- **Increase trustworthiness of prediction**

This is because the columns that are removed that provide information about the target information would influence eLCS ability to learn and accurately predict the target without genuinely learning from the other features.

4. Improved LCS-Based System

☐ Students must improve the LCS code, LCS configuration, or machine learning pipeline to develop a novel improved system. Possible improvements include:

- ☐ Hyperparameter tuning.
- ☐ Improved preprocessing designed for LCS performance.
- ☐ Feature selection or dimensionality reduction.
- ☐ Rule pruning or rule compaction.
- ☐ Modified LCS algorithm.

- ☐ Class-imbalance handling.
- ☐ Cost-sensitive learning.
- ☐ Students are required to:
 - ☐ Clearly identify what was changed compared with the original LCS system.
 - ☐ Justify why the change may improve prediction accuracy or model quality.
 - ☐ Provide a brief algorithmic explanation, pseudocode, or flowchart of the improved system.
 - ☐ Compare the improved system with the original LCS baseline.

4.1 Summary of changes vs. the original LCS system

4.2 Justification – expected effect on accuracy or model quality

4.3 Algorithmic explanation (pseudocode or flowchart)

4.4 Comparison against the original LCS baseline

5. Experiments

- ☐ Students must use a suitable and justified validation strategy. Students are required to:
 - ☐ Use either an 80/20 train-test split or k-fold cross-validation.
 - ☐ Justify the selected validation approach.
- ☐ Students must evaluate results using appropriate metrics and statistical tests. Students are required to:
 - ☐ Select evaluation metrics suitable for the selected problem.
 - ☐ For classification, consider accuracy, precision, recall, F1-score, balanced accuracy, ROC-AUC, PR-AUC, and confusion matrix.
 - ☐ For regression, consider MAE, MSE, RMSE, MAPE, R^2 , and residual analysis.
 - ☐ Justify the selected metrics.
 - ☐ Apply at least one suitable statistical test to compare model performance.
 - ☐ Interpret whether performance differences are statistically significant.

5.1 Validation strategy (80/20 split or k-fold CV) – justification

5.2 Evaluation metrics – selection and justification (classification: accuracy, precision, recall, F1, balanced accuracy, ROC-AUC, PR-AUC, confusion matrix; regression: MAE, MSE, RMSE, MAPE, R^2 , residual analysis)

5.3 Statistical significance testing

5.4 Interpretation of significance results

6. Model Comparison

- ☐ Students must compare the proposed improved LCS-based system with the original LCS and at least three conventional machine learning models. The comparison must include:
 - ☐ Original LCS on the raw or minimally processed dataset.
 - ☐ Original LCS on the preprocessed dataset.
 - ☐ Improved LCS-based system on the preprocessed and/or feature-engineered dataset.
 - ☐ At least three conventional non-deep-learning machine learning models.
- ☐ Conventional models may include Random Forest, Support Vector Machine, Decision Tree, Logistic Regression, Naive Bayes, k-Nearest Neighbours, Gaussian Process, or other appropriate non-deep-learning models. Deep-learning models are not permitted for the required model comparison.

6.1 Original LCS – raw/minimally processed data

6.2 Original LCS – preprocessed data

6.3 Improved LCS – preprocessed/feature-engineered data

6.4 At least three conventional ML models (e.g. Random Forest, SVM, Decision Tree, Logistic Regression, Naive Bayes, k-NN, Gaussian Process) – no deep learning

6.5 Comparative results table

7. Interpretation of Results and Explainable AI

- ☐ Students must include a separate section on result interpretation and Explainable AI. Students are required to:
 - ☐ Interpret the main findings from the model comparison.
 - ☐ Explain important LCS rules, feature conditions, or classifier patterns.
 - ☐ Use at least three sample LCS rules/classifiers to discuss the decision-making process.
 - ☐ Discuss how LCS rule-based outputs support interpretability.
 - ☐ Discuss whether the explanations are trustworthy and practically meaningful.

7.1 Interpretation of main findings from the model comparison

7.2 Explanation of key LCS rules / feature conditions / classifier patterns

7.3 At least three sample LCS rules discussed for decision-making insight

7.4 How rule-based outputs support interpretability

7.5 Trustworthiness and practical meaningfulness of the explanations

8. Discussion, Limitations, and Future Work

☐ Students must critically discuss the results and the responsible use of the developed system. Students are required to discuss:

- ☐ Reasons for performance improvement or lack of improvement.
- ☐ Practical meaning of the results, not only numerical performance.
- ☐ Limitations due to data quality, dataset size, class imbalance, noise, missingness, or feature relevance.
- ☐ Future work, such as improved LCS variants, better preprocessing, external validation, or domain expert review.

8.1 Reasons for performance improvement (or lack thereof)

8.2 Practical significance of results beyond raw numbers

8.3 Limitations – data quality, dataset size, class imbalance, noise, missingness, feature relevance

8.4 Future work – improved LCS variants, better preprocessing, external validation, domain expert review

Notes

Final report: 5,000--6,500 words, excluding references and appendices.

Source code or Jupyter Notebook.

GitHub repository link and commit-history evidence.

Contribution statement, including a brief individual reflection from each group member.

Cleaned and feature-engineered dataset, or clear dataset reproduction instructions.

Turnitin similarity report with a similarity score not exceeding 20%.

References

Dataset

- <https://www.kaggle.com/datasets/danishjmeo/281k-patient-dataset-for-multi-disease-prediction?resource=download>

Sub Datasets

- <https://www.kaggle.com/datasets/iammustafatz/diabetes-prediction-dataset>,
- [kaggle.com/datasets/oktayrdeki/heart-disease](https://www.kaggle.com/datasets/oktayrdeki/heart-disease),
- [kaggle.com/datasets/ankushpanday1/hypertension-risk-prediction-dataset](https://www.kaggle.com/datasets/ankushpanday1/hypertension-risk-prediction-dataset)

Research

Latest heart disease statistics. (n.d.). Heart Foundation NZ.

<https://www.heartfoundation.org.nz/statistics>